

SENPai: Self-Experimentation for Physical AI

An Observability-Based Research Harness

Thomas Capelle^{*12} Morgan McGuire^{*12} Justin Hodges²




Figure 1. Results of 167 experiments from the SENPAI experiment ledger for the five TandemFoilSet-Balanced trials, grouped by count of experiments per trial, the experiment hypothesis family and final PR outcome depending on if it was merged, closed or left unresolved at the 12-hour trial cut-off.

Abstract

SENPai (Self-Experimentation for Physical AI) is an observability-first research harness in which multi-agent state is grounded in pull requests and structured experiment logs rather than agent memory and scratchpads. Although task-agnostic, we evaluate SENPAI on CFD-surrogate recipe search, where performance depends on architecture, optimization, losses, normalization, and physics-aware considerations. SENPAI uses a thin Advisor/Student loop to turn hypotheses into PRs, training runs, and PR comments, producing an experiment ledger queryable by both agents and researchers. This makes experiments auditable, recoverable, and researcher-steerable. We evaluate SENPAI across DrivAerML, Air-

FRANS, and TandemFoilSet. Starting from the original Transolver model, SENPAI obtains the best single-model DrivAerML surface- and volume-pressure relative- L_2 errors among compared references while improving wall-shear-stress error, reduces AirfRANS surface-MSE below all compared references, and reaches lower TandemFoilSet normalized full-field MSE on the cruise-random-uniform split than the reported benchmark. Overall, SENPAI supports sparsely steered, multi-day semi-autonomous research that can deliver strong task-specific models. We release the harness¹ and Weights & Biases experiment ledger.²

^{*}Equal contribution ¹Weights & Biases, San Francisco, CA, USA ²CoreWeave, Roseland, NJ, USA. Correspondence to: Morgan McGuire <morg@wandb.com>.

¹<https://github.com/wandb/senpai>

²See Appendix B.

1. Introduction

CFD surrogates offer a way to accelerate engineering design: learned models can screen early-stage geometries, explore operating regimes, and focus expensive CFD or physical testing on the most promising candidates. Recent surrogate models now handle larger meshes, irregular geometries, stricter accuracy targets, and smaller training sets than were feasible only a few years ago (Zhou et al., 2026), making hybrid ML-CFD workflows increasingly attractive.

In practice, however, training a useful CFD surrogate is still a niche capability. Performance depends on architecture, optimization, normalization, boundary conditions, symmetries, and rollout stability (Wang et al., 2024). Strong recipes require both ML expertise and domain knowledge in aerodynamics or fluid dynamics (Hodges, 2024), which many engineering teams do not have. SENPAI targets this gap by exploring, documenting, and refining problem-specific surrogate recipes while keeping researchers in the loop.

Recent autonomous-research systems show that LLM agents can search literature, propose hypotheses, edit code, run experiments, and produce reports (Lu et al., 2026; Schmidgall et al., 2025; Yamada et al., 2025; Miyai et al., 2026). SENPAI builds on this direction but focuses on a different systems question: how should long-running, multi-experiment ML research be represented so that both agents and researchers can audit, recover, compare, and steer it?

SENPAI’s central design choice is to ground the research loop in external systems of record that ML teams already inspect: pull requests and git history for hypotheses, code changes, review, and provenance (GitHub Docs, 2026; Chacon & Straub, 2014); and a hosted experiment tracker for configurations, metrics, checkpoints, system telemetry, and agent traces (Biewald, 2020; Weights & Biases, 2026). Our implementation uses GitHub and Weights & Biases, but the harness design is not specific to either service. By using these tools as an experiment ledger, each hypothesis, code diff, training result and review decision becomes both machine-queryable and researcher-reviewable.

The harness loop is deliberately thin. An Advisor reads the current experiment ledger, proposes a literature-grounded hypothesis as a draft PR, and assigns it to a GPU-backed Student. The Student checks out the branch, modifies the training code, runs the experiment, logs metrics to the experiment tracker, and reports results in the PR. The Advisor then merges, requests revision, or closes the line of inquiry. Humans can steer the Advisor or Student through GitHub issues and PR comments, while markdown scratchpads remain secondary sources of state rather than the primary system of record.

Across the research programme, SENPAI produced over

3,000 PRs, 11,000 experiment-tracker runs, and operated with up to 59 concurrent Student agents, turning a multi-day semi-autonomous research programme into a durable experiment ledger. Figure 1 shows one compact slice of this ledger: 167 TandemFoilSet-Balanced PRs grouped by research round, hypothesis family, and final outcome.

Contributions.

- **Observable research state.** We introduce SENPAI, a semi-autonomous ML research harness that grounds experiment state in PRs, commits, and experiment-tracker runs, producing a ledger that is queryable by agents and inspectable by researchers.
- **CFD recipe search.** Starting from Transolver-style baselines, SENPAI discovers improved recipes across AirFRANS, DrivAerML, and TandemFoilSet, including best single-model DrivAerML pressure results among compared references.
- **Failure modes at scale.** We analyze a 24-hour fleet trace with 53,022 Claude requests and 5.24B tokens, showing that the dominant failures were monitor-driven context bloat, brittle tool interfaces, checkpoint/result-capture gaps, and state reconciliation issues rather than failures of scientific reasoning.

2. Related Work

Recent autonomous-research systems show that LLM agents can search literature, propose hypotheses, edit code, run experiments, and produce research reports. The AI Scientist generates ideas, implements experiments, writes papers, and runs automated review (Lu et al., 2026; Sakana AI, 2024), while AI Scientist-v2 extends this direction with agentic tree search and workshop-level paper generation (Yamada et al., 2025; Sakana AI, 2025). Agent Laboratory structures literature review, experimentation, and report writing around a researcher-provided idea (Schmidgall et al., 2025), and Jr. AI Scientist iteratively improves a baseline paper using modern coding agents (Miyai et al., 2026). Related systems emphasize collaboration, shared artifacts, or domain tools: ResearchAgent generates and revises ideas over scientific-literature graphs (Baek et al., 2025), AgentRxiv lets agent laboratories upload and retrieve generated reports from a shared preprint server (Schmidgall & Moor, 2025), ChemCrow and Coscientist connect LLM planners to chemistry tools and laboratory workflows (Bran et al., 2024; Boiko et al., 2023), and FunSearch couples an LLM generator to a verifiable evaluator and scored program database (Romera-Paredes et al., 2024). Similarly, in general ML engineering, Hugging Face’s ml-intern is a recent, non-science-focused example that reviews papers, writes training code, and runs experiments (Hugging Face, 2026). These systems establish the value of tool use, execu-

tion, and persistent artifacts in autonomous research. SENPAI builds on that direction, but studies a different substrate question: how should the state of a long-running, multi-experiment ML research programme be represented so that both agents and researchers can inspect, compare, recover, and steer it?

The closest systems to SENPAI also recognize that long-horizon agents need state outside the model context, but they choose different substrates. Kosmos maintains an internal structured memory of literature-search and data-analysis summaries to choose later tasks, and cites final report claims with papers or generated Jupyter notebooks (Mitchener et al., 2025). AiScientist uses a permission-scoped File-as-Bus workspace so agents can re-ground on file-based analyses, plans, code, logs, and experimental evidence (Chen et al., 2026). Grounded autonomous research externalizes state as rerunnable simulation artifacts and comparison results within a single-paper reproduction loop (Huang, 2026). SENPAI takes a complementary systems position: the authoritative state should be the MLOps tooling machine learning researchers already use. In SENPAI, pull requests carry hypotheses, code review, discussion, and review decisions (GitHub Docs, 2026); git history provides ordered code provenance (Chacon & Straub, 2014); and hosted experiment trackers provide structured records of metrics, hyperparameters, artifacts, checkpoints, system telemetry, and traces (Biewald, 2020; Weights & Biases, 2026). This distinction matters most at scale: once an autonomous research system produces hundreds or thousands of experiments, the central challenge is no longer only whether an agent can run an effective experiment, but whether researchers can understand the programme it is creating.

3. System Design

Core principle. SENPAI’s central design principle is that authoritative experiment state lives in pull requests, git history, and experiment-tracker runs, not in agent memory or local scratchpads. The harness is therefore organized around making each hypothesis, code change, training result, and review decision durable, queryable, and inspectable through existing research tools.

PR-mediated lifecycle. SENPAI is a thin two-role harness over a task repository. An Advisor reads the current experiment ledger, optionally calls a literature-search sub-agent, opens draft PRs containing hypotheses and experiment instructions, and assigns them to GPU-backed Students. Each Student claims its labeled PR, edits the branch, launches training with experiment-tracker logging, and reports results back as PR comments. Researchers can steer Advisors or Students via labeled GitHub issues or by leav-




Figure 2. SENPAI deployment architecture. A thin harness deploys one Advisor and N Student pods; both roles can call a shared literature-search sub-agent. All programme state lives outside the cluster in the experiment ledger: PRs, git history, and experiment-tracker runs.

ing comments on a PR. A hypothesis’s trajectory therefore remains visible as a standard PR history linked to experiment-tracker runs, rather than as an agent-internal trace.

Design choices. SENPAI follows three design choices: externalizing experiment state into existing systems of record, keeping the autonomous loop thin, and exposing the same ledger to both agents and researchers. The role prompts, problem-program contract, helper skills, and outer-loop pseudocode are summarized in Appendices A and D to F.

3.1. Externalized Experiment State

The experiment ledger has two linked records. The PR records the hypothesis, implementation instructions, code diff, Student report, Advisor review, and final decision. The experiment tracker records the run details: configuration, metrics, checkpoints, system data, and agent traces. In our experiments, this tracker was Weights & Biases. Together these records are inspectable through familiar code-review and experiment-tracking interfaces and queryable through mature CLIs, allowing both researchers and agents to audit individual experiments or summarize programme-level progress.

The Advisor also maintains lightweight markdown summaries for focus, such as current baseline metrics, dataset notes, and open research directions. These files are treated as caches rather than authoritative state: when summaries disagree with PRs, git history, or experiment-tracker runs, the ledger wins.

3.2. Thin Advisor/Student Loop

Each experiment follows a PR-mediated lifecycle. The Advisor drafts a PR with a hypothesis, implementation guid-

ance, target metrics, baseline values, and a Student label. The Student implements the change on the PR branch, runs training with experiment-tracker logging, and posts commands, training run links, metrics, and experiment interpretation as a PR comment. Once the PR is marked ready for review, the Advisor compares the result against the stated baseline and either merges the improvement, requests revision, or closes the branch. Merged PRs advance the programme baseline used for subsequent hypotheses.

SENPAI deploys as a small set of Kubernetes workloads: one Advisor CPU pod, N Student GPU pods, and a shared persistent data volume for datasets and checkpoints. The harness image is stable across problems; a task repository can be swapped in per SENPAI research programme, so every agent commit, branch, and PR lives in a self-contained target repository that a researcher can inspect independent of the harness.

The outer harness loop is programmatic rather than agentic. A shell wrapper polls GitHub and Student state for review-ready PRs, human issues, new comments, and idle Students; if no action is available, it sleeps without invoking the LLM. When action is required, the wrapper starts or resumes the Advisor with a triage message. This keeps routine polling outside the model context and focuses the agent on research decisions rather than heartbeat bookkeeping.

3.3. Human Steering Through Shared Observability

Researchers primarily steer SENPAI’s Advisor or Students through GitHub issues, while PR comments can also provide experiment-level intervention. This channel is intentionally low bandwidth: SENPAI is designed for sparse steering rather than continuous supervision, and each intervention remains attached to the public experiment record it affected.

In our longest observed SENPAI deployment on DrivAerML, the system ran continuously for 9 days under a 4.5-hour per-experiment budget, opening 289 PRs across an 8-student active fleet with 10 issue-based interventions from researchers.

4. Experiments and Results

4.1. Experiment Setup

Harness. Claude Code (v2.1.85 to v2.1.117) was run in headless mode as the agent harness for both the Advisor and Students. Claude Opus 4.6 and Claude Opus 4.7 were used with a compaction trigger at 200k tokens. Student pods were scheduled on NVIDIA RTX 6000 Blackwell nodes with 96 GB VRAM per GPU, using from 1 to 8 GPUs per Student node depending on the experiment setting. At maximum, $N = 59$ concurrent Students were tested, with peak

concurrent training runs of 347. Over the five-week experimentation and harness-development cycle, 30 billion tokens were processed by Claude Code while running physical-AI experiments.

Scale. Over the course of the research programme, SENPAI created over 3,000 PRs and recorded over 11,000 experiment-tracker runs while deploying, at peak, 59 distinct Student agents. Approximately 30 billion Claude Code tokens were processed during experimentation.

4.2. Datasets and Metrics

To demonstrate SENPAI’s CFD performance across multiple benchmarks, we train and evaluate on three CFD surrogate datasets spanning 3D automotive aerodynamics, 2D airfoil RANS, and 2D tandem-airfoil flow: DrivAerML (Ashton et al., 2024), AirfRANS (Bonnet et al., 2022), and TandemFoilSet (Lim et al., 2026). Figure 3 shows representative geometries from all three datasets.

DrivAerML. We use the public processed-case split, 400/34/50 train/validation/test, and report surface-pressure, vector wall-shear-stress, and volume-pressure relative- L_2 , computed per case on unnormalized predictions and then averaged over test cases. Each raw DrivAerML case contains approximately 8.8M surface points and 160M volume points.

AirfRANS. We use the official full task with the last 10% of the official training list held out for validation, giving a 720/80/200 train/validation/test split while keeping the official test set unchanged. Each simulation has roughly 150K points. We test on normalized targets on the official test split.

TandemFoilSet. We use the paper’s original Experiment 4 partition, Cruise Random and Race Car tasks, as one of two TandemFoilSet datasets and measure denormalized surface-pressure MAE. TandemFoilSet cases average approximately 187K points; the active TandemFoilSet-Balanced subset averages approximately 141K points. In addition, we introduce a second dataset split, TandemFoilSet-Balanced (McGuire & Capelle, 2026), with 1,499 training cases across four balanced test splits: one single-foil in-distribution split, two unseen-front-camber geometry holdouts (race-car and cruise), and one Reynolds-number holdout. An equal-weight average denormalized surface-pressure MAE of these four splits is used as the test metric.

4.3. DrivAerML

Result. On the DrivAerML benchmark, SENPAI’s best single trained checkpoint is H147 from PR #1344, W&B




Figure 3. CFD surrogate problems used by SENPAI, shown as representative dataset geometries rather than benchmark results. (a) DrivAerML is a 3D automotive aerodynamics task. (b) AirfRANS contains families of 2D airfoil cross-sections. (c) TandemFoilSet contains 2D tandem-airfoil cases; the panel shows four geometry-diverse tandem cases, with two foils per case, drawn from race-car and cruise configurations.

Table 1. DrivAerML test relative- L_2 error (%). Lower is better.

Method	p_s	τ	p_v
SENPAI (ours)	3.56	6.54	3.40
Transolver-3 (Zhou et al., 2026)	3.71	5.85	5.72
AB-UPT (Alkin et al., 2025)	3.82	7.29	6.08
Transolver++ (Luo et al., 2025)	4.12	6.42	6.70
Transolver (Wu et al., 2024)	4.81	8.95	7.74

run `k6q4c3on`. Starting from a Transolver-based implementation, this checkpoint improves on the original Transolver for surface-pressure, vector wall-shear-stress, and volume-pressure relative- L_2 error. It also gives the best single-model pressure results among the public references in Table 1: 3.5634% surface pressure versus 3.71% for Transolver-3, and 3.4014% volume pressure versus 5.72% for Transolver-3. Wall-shear error is 6.5409%, improving on AB-UPT and Transolver while remaining behind Transolver-3 and Transolver++. The full recipe, run provenance, and per-axis wall-shear diagnostics are in Appendix G. Figure 4 shows the semi-autonomous experiment trajectory for surface-pressure test error across earlier SENPAI campaigns.

Recipe. The selected single-model recipe SENPAI found used the original Transolver backbone but scaled it to 6 layers, 512 hidden units, and 128 slices, replaced fixed sinusoidal coordinates with a learnable Fourier positional embedding initialized across multiple spatial scales (Li et al., 2021; Schenck et al., 2025), and trained with Lion (Chen et al., 2023), EMA, GradNorm adaptive loss balancing (Chen et al., 2018), curvature attention bias, and y -

Table 2. AirfRANS full-task test MSE. Targets are train-statistic normalized, per-case means on the official test split; lower is better. Spider denotes SpiderSolver (Qi et al., 2025).

Metric	SENPAI	Spider	GeoANF	Transolver
Surface MSE	0.00130	0.0043	0.0089	0.0142
Volume MSE	0.00451	0.0017	0.0062	0.0037

symmetry augmentation. It trained with 8-GPU DDP for an effective batch size of 8 under a 30-epoch budget, with the best EMA checkpoint selected at epoch 25. A notable empirical finding from SENPAI was the application of this lightweight STRING-inspired encoding from computer vision and robotics use cases to CFD surrogate modelling.

4.4. AirfRANS

Result. On the AirfRANS full benchmark, SENPAI, starting from a Transolver-based implementation, improves substantially on the original Transolver surface-MSE result (5-seed mean 0.0013 versus 0.0142) and on the strongest published surface reference, SpiderSolver (0.0043). The same model obtains mean volume-MSE of 0.00451, improving on GeoANF (Xiao et al., 2024) but remaining above SpiderSolver and slightly above the original Transolver volume result. SENPAI therefore achieves the strongest reported surface-MSE among these references, while the full benchmark remains limited by volume accuracy.

Recipe and caveat. The final AirfRANS recipe SENPAI found used a 2-layer, 256-hidden-unit, 4-head SENPAI-Transolver with 96 slices, fixed multiscale Fourier coor-




Figure 4. Evolution of DrivAerML surface-pressure test error across two SENPAI campaigns. Faded markers show individual completed experiments from the `yi` and `tay` branches, while solid curves show the cumulative Pareto frontier: the best surface-pressure relative L_2 error achieved up to each experiment count. The vertical lines mark GitHub-recorded human researcher/operator interventions. `tay` was instructed to use DDP8 and was given a 6 hour per-experiment training budget, hence the better performance.

dinate features, a zero-initialized surface refinement head, Lookahead-wrapped AdamW, cosine scheduling, gradient clipping, weight decay, EMA with target decay 0.999, and a volume-weighted normalized-MSE objective with volume loss weight 10.0. All five seeds beat the published SpiderSolver surface-MSE result, while volume-MSE remains above SpiderSolver and more sensitive to seed. The complete five-seed configuration is reported in Appendix H.

4.5. TandemFoilSet

Result. Under the TandemFoilSet Experiment-4 task, SENPAI found a Transolver-based improvement on the cruise random uniform full-field MSE of the original paper ($1.7 \cdot 10^{-3}$ versus $1.0 \cdot 10^{-1}$). On TandemFoilSet-Balanced, SENPAI found a recipe that produced a denormalized surface-pressure MAE of 23.45 using a 5-seed average.

Recipe signal. For this benchmark, SENPAI used learning rates of $\sim 1e-4$, gradient clipping, EMA, warmup, and schedule-length tuning.

Table 3. TandemFoilSet-Balanced 5-seed mean denormalized surface-pressure MAE. Lower is better.

Metric	MAE
Single-foil in-distribution	17.2905
Race-car camber holdout	31.2177
Cruise camber holdout	28.3145
Reynolds-number holdout	16.9834
Average	23.4515

4.6. Flat-Agent Comparison

Setup. To contrast SENPAI’s Advisor-mediated loop, we compared the performance of a `kagent` harness (Capelle & McGuire, 2026) against the SENPAI harness using the TandemFoilSet-Balanced dataset and the same compute budget. `kagent` is a lighter-weight harness that drops the Advisor/Student split in favor of a flat Kaggle-like competition with a shared leaderboard; each agent keeps its own experiment journal and is prompted to compete aggressively to climb the leaderboard. In contrast to a Kaggle competition, the agents also have access to the same hosted experiment-logging project and commit history, so they can learn from other agents’ work. The `kagent` prompt and the full run summary are provided in Appendix I.

Table 4. TandemFoilSet-Balanced 12-hour harness comparison. Lower average surface-pressure MAE is better.

Metric	SENPAI	kagent
Best avg-surf- p MAE	44.15	34.58

Interpretation. In a direct 12-hour head-to-head, kagent behaved like a public Kaggle sprint: agents trained short variants, watched the leaderboard and shared hosted-logger configs, and quickly reused recipes that worked for peers. This fast peer diffusion, including copied hyperparameters and occasional prediction ensembles, likely gave it an exploitation-speed edge over SENPAI’s slower PR-mediated review loop. While this was successful in a 12-hour experiment setting, the inability to steer these agents and the less rigorous polling and monitoring infrastructure are likely to lead to drift and context rot as each Claude Code agent’s context window fills. However a later external validation potentially contradicts this. A model checkpoint from a 30-hour autonomous kagent run, with a 30 minute per-experiment budget, was submitted to the GRaM ICLR 2026 competition and subsequently ranked 4th of 22 entries on the official held-out leaderboard, with relative- L_2 error 0.0553 ± 0.0168 .

4.7. Modded-NanoGPT Track 3

Result. To demonstrate SENPAI’s applicability beyond CFD research, we also ran SENPAI on Keller Jordan’s modded-NanoGPT Track 3 optimization benchmark (Jordan, 2026). SENPAI produced a candidate Track 3 result at 2925 steps, state-of-the-art as of writing, and pending official ratification. The result is recorded in PR #<BEST_RESULT_PR>. We estimate that one average SENPAI trial corresponds to roughly 3.5 zettaFLOP of hardware-peak training compute; see Appendix C for the training recipe and further details.

5. Discussion

5.1. Observable State as a Control Surface

Control surface. The PR and experiment-tracker ledger made SENPAI useful as a research assistant rather than a tool to execute experiments. Because hypotheses, code diffs, metrics, failures, and review decisions were all recorded in systems researchers already use, humans could inspect the research programme without needing to read the agents’ context. The same records also gave agents a stable system of record: they could query prior runs, identify stale experiments, compare seed behavior, and decide whether a new result actually advanced the baseline. In SENPAI, observability is therefore not simply logging; it is the layer through which both humans and agents can steer the re-

search programme.

Recovery. This externalized state also reduced the fragility of long-running agent sessions. When agents compacted, restarted, or lost local scratchpad context, the next cycle could reconstruct the experiment from the PR description, code diff, git history, experiment-tracker records, and posted result comments. Local summaries remained useful for focus, but they were treated as caches over a more durable experiment record.

Research programme. The ledger also turned a collection of individual experiments into a queryable research programme. Researchers and agents could ask what had been tried, which failures were informative, which recipes transferred across datasets, where progress had stalled, and which resources were idle. This is SENPAI’s central systems claim: the goal is not to replace the researcher, but to increase the bandwidth at which researchers can understand, guide, and accelerate experimentation.

5.2. Harness Tradeoffs

Tradeoff. The kagent comparison highlights a harness-design tradeoff. Flat, leaderboard-driven cohorts can share successful recipes quickly because agents observe peer results directly and optimize for short-horizon score improvement. In contrast, SENPAI’s PR-mediated loop is slower but preserves review decisions, experiment provenance, and sparse human steering. We view kagent as useful for rapid exploitation, and SENPAI as better suited to longer-running programmes where auditability and recovery matter alongside score.

5.3. Failure Modes: Where Long-Running Research Agents Break

Observed failures. SENPAI’s most common failures were harness-engineering failures: monitoring long-running jobs, recovering across compaction or restarts, enforcing brittle tool interfaces, and reconciling PR state with experiment-tracker state. They were less often failures of scientific reasoning. This distinction matters because it suggests that improving autonomous research systems requires better observable infrastructure, not only stronger prompts or more capable models. Additional artifact schemas and the extended taxonomy are listed in Appendix J.

Monitor bloat. During development, one dominant cost to running SENPAI was monitor-driven context bloat. In one notable experiment, Students monitored training logs for progress, errors, checkpoints, and completion, but persistent `tail -f | grep` events repeatedly resumed long-lived Student sessions for low-information log lines. In a 24-

hour window with 57 Students running, 615 monitor setups produced 28,247 agent-facing monitor events and 28,246 agent responses, consuming 3.25B cache-inclusive tokens. The median agent response reloaded roughly 110k cached tokens while adding only 347 new uncached input tokens. This is an example of an architectural failure mode when running long-running research systems: training monitors should be cheap and escalation-based, surfacing only relevant events such as errors, experiment completion, or metric milestones.

Tool interfaces. Tool-use errors were smaller but systematic: in the same 24-hour window, 892 of 25,365 tool results had errors. The main causes were failed shell commands, GitHub scope errors, blocked wait patterns, wrong paths, failed pushes, oversized reads, and edit-guard failures. These errors point to a practical lesson: high-frequency harness operations should be executable helpers with narrow interfaces, rather than repeatedly reconstructed shell or GitHub operations inside agent context.

State drift. Compaction introduced a separate risk. The 24-hour window contained 240 automatic compactions across the 57 Students, with Student contexts reduced from roughly 196k tokens to 3k-token summaries. These summaries were sufficient for resumption, but lossy: state drift appeared in the agent conversation logs with mistakes such as incorrect GitHub labels or confusion over live experiment state. SENPAI’s experiment ledger mitigated this risk by externalizing much of the key experiment state. As a result, in that 24-hour window, despite an average of 4.2 compactions per Student, 38 of 39 Student sessions whose monitors showed training had completed also left PR comments as expected, and 33 of 39 PRs reached ready/status-review handoff as expected, indicating that despite multiple compactions, Students were able to use the experiment ledger’s external state to course-correct and take experiments to completion.

Mitigation. From observing these failure modes, we conclude that long-running research agents need robust, observable infrastructure more than additional prompt prose. Monitor bloat requires thoughtful observation; tool errors require executable helpers; training failures require mandatory result and checkpoint capture; state drift requires a single authoritative ledger, with summaries treated as caches. SENPAI did not eliminate failure, but it made failure measurable, attributable, and repairable.

5.4. Limitations

Scope. SENPAI is task-agnostic as a harness, but our empirical evidence comes from CFD-surrogate recipe search. We have not yet evaluated whether the same design transfers

to domains with different data modalities, evaluation loops, or laboratory constraints.

Surface metrics. The strongest empirical results are on surface quantities. This matches many aerodynamic design workflows, where surface pressure and forces are primary signals. Volume-field prediction remains less developed on AirfRANS, where SENPAI improves surface MSE while remaining behind the strongest published volume-MSE method; on DrivAerML, the selected single checkpoint improves both pressure fields while wall-shear remains behind the strongest public wall-shear reference.

Supervision. SENPAI is semi-autonomous rather than fully autonomous. It requires an initial task repository, metric definition, dataset setup, and occasional human steering. The GitHub Issues and PR-comment interface is auditable and low bandwidth, but slower than interactive supervision.

Infrastructure dependence. Finally, SENPAI depends on mature external systems of record and substantial compute/API budget. In our implementation these were GitHub and Weights & Biases, which improved observability but also introduced dependence on hosted-service reliability, access control, rate limits, and retention policies. We release the harness and experiment ledger so future work can audit costs, reproduce results, and evaluate more efficient variants.

6. Conclusion

SENPAI shows that a lightweight, observability-first harness can sustain semi-autonomous ML research when multi-agent state is grounded in the artifacts researchers already inspect: pull requests, git history, and experiment-tracker runs. Across a multi-day CFD-surrogate research programme, this design supported thousands of autonomous experiments over three benchmarks while allowing researchers to steer the system through ordinary PR review and a GitHub Issues channel.

The main lesson is that reliability comes from the harness, not the agent instructions alone. Long-running research agents did not primarily fail because they lacked clear instructions; they failed at system boundaries: brittle tool interfaces, unbounded monitoring, missing checkpoint capture, and drift between derived summaries and the authoritative ledger. Prompt hardening alone did not remove these failure modes. Durable external state, clear scripts for passing work between agents, filtered monitoring, and mandatory result capture were more important for keeping the research programme recoverable and auditable.

SENPAI’s contribution is therefore not to replace the researcher, but to increase the bandwidth at which researchers

can understand, guide, and accelerate autonomous experimentation. By making both the inner experiment loop and the outer harness loop observable through mature systems of record, SENPAI turns autonomous research from an opaque agent workflow into a queryable, reviewable research programme. We release the harness, experiment ledger, and failure-mode analysis to support future work on more reliable and efficient autonomous research systems.

Software and Data

The SENPAI harness and kagent comparator are released publicly. Public experiment ledgers are available through GitHub PR histories in [wandb/senpai](#) and [TandemFoilSet-Balanced](#), with W&B run records under the projects listed in Appendix B.

Impact Statement

This work aims to make autonomous ML research more auditable, recoverable, and steerable by grounding agent activity in systems of record that researchers already inspect. The positive impact is increased researcher bandwidth for expensive scientific ML experimentation, especially in domains where strong baselines require repeated empirical iteration. The main risks are increased compute/API consumption, over-trusting autonomous experiment choices, and propagating errors when agent-written summaries drift from measured evidence. SENPAI mitigates these risks by keeping humans in the review loop, preserving experiment provenance in PRs and tracked runs, and treating summaries as caches rather than authoritative state.

References

- Alkin, B., Bleeker, M., Kurle, R., Kronlachner, T., Sonnleitner, R., Dorfer, M., and Brandstetter, J. AB-UPT: Scaling neural CFD surrogates for high-fidelity automotive aerodynamics simulations via anchored-branched universal physics transformers. *Transactions on Machine Learning Research*, 2025. ISSN 2835-8856. URL <https://openreview.net/forum?id=nwQ8nit1TZ>.
- Ashton, N., Mockett, C., Fuchs, M., Fliessbach, L., Hetmann, H., Knacke, T., Schonwald, N., Skaperdas, V., Fotiadis, G., Walle, A., Hupertz, B., and Maddix, D. DrivAerML: High-fidelity computational fluid dynamics dataset for road-car external aerodynamics, 2024. URL <https://arxiv.org/abs/2408.11969>.
- Baek, J., Jauhar, S. K., Cucerzan, S., and Hwang, S. J. ResearchAgent: Iterative research idea generation over scientific literature with large language models. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 6709–6738, Albuquerque, New Mexico, Apr 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.naacl-long.342. URL <https://aclanthology.org/2025.naacl-long.342/>.
- Biewald, L. Experiment tracking with Weights and Biases, 2020. URL <https://www.wandb.com/>. Software available from wandb.com.
- Boiko, D. A., MacKnight, R., Kline, B., and Gomes, G. Autonomous chemical research with large language models. *Nature*, 624:570–578, 2023. doi: 10.1038/s41586-023-06792-0. URL <https://www.nature.com/articles/s41586-023-06792-0>.
- Bonnet, F., Mazari, J., Cinnella, P., and Gallinari, P. AirfRANS: High fidelity computational fluid dynamics dataset for approximating Reynolds-Averaged Navier–Stokes solutions. In *Advances in Neural Information Processing Systems*, volume 35, 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/94ab7b23a345f93333eac8748a66c763-Abstract-Datasets_and_Benchmarks.html.
- Bran, A. M., Cox, S., Schilter, O., Baldassari, C., White, A. D., and Schwaller, P. Augmenting large language models with chemistry tools. *Nature Machine Intelligence*, 6: 525–535, 2024. doi: 10.1038/s42256-024-00832-8. URL <https://www.nature.com/articles/s42256-024-00832-8>.
- Capelle, T. and McGuire, M. kagent: A lightweight multi-agent experiment harness. GitHub repository, 2026. URL <https://github.com/tcapelle/kagent>. Software artifact.
- Chacon, S. and Straub, B. *Pro Git*. Apress, 2 edition, 2014. ISBN 978-1-4842-0076-6. doi: 10.1007/978-1-4842-0076-6. URL <https://git-scm.com/book/en/v2>.
- Chen, G., Chen, J., Chen, L., Zhao, J., Meng, F., Zhao, W. X., Song, R., Chen, C., Wen, J.-R., and Jia, K. Toward autonomous long-horizon engineering for ML research, 2026. URL <https://arxiv.org/abs/2604.13018>.
- Chen, X., Liang, C., Huang, D., Real, E., Wang, K., Liu, Y., Pham, H., Dong, X., Luong, T., Hsieh, C.-J., Lu, Y., and Le, Q. V. Symbolic discovery of optimization algorithms. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/9a39b4925e35cf447ccba8757137d84f-Abstract-Conference.html.
- Chen, Z., Badrinarayanan, V., Lee, C.-Y., and Rabinovich, A. GradNorm: Gradient normalization for adaptive loss

- balancing in deep multitask networks. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pp. 794–803. PMLR, 2018. URL <https://proceedings.mlr.press/v80/chen18a.html>.
- GitHub Docs. About pull requests, 2026. URL <https://docs.github.com/en/pull-requests/collaborating-with-pull-requests/proposing-changes-to-your-work-with-pull-requests/about-pull-requests>. Accessed 2026-04-24.
- Hodges, J. *Approaching Machine Learning Problems in Computational Fluid Dynamics and Computer Aided Engineering Applications: A Monograph for Beginners*. Justin Hodges, 2024. ISBN 9798878702317.
- Huang, H. Towards grounded autonomous research: An end-to-end LLM mini research loop on published computational physics, 2026. URL <https://arxiv.org/abs/2604.12198>.
- Hugging Face. ml-intern: An open-source ML engineer that reads papers, trains models, and ships ML models, 2026. URL <https://github.com/huggingface/ml-intern>.
- Jordan, K. modded-nanoGPT: Optimization benchmark. GitHub repository, 2026. URL https://github.com/KellerJordan/modded-nanogpt/tree/master/records/track_3_optimization. Track 3 benchmark; accessed 2026-05-29.
- Karpathy, A. Reproducing GPT-2 (124m) in llm.c. GitHub discussion, 2024. URL <https://github.com/karpathy/llm.c/discussions/481>.
- Li, Y., Si, S., Li, G., Hsieh, C.-J., and Bengio, S. Learnable Fourier features for multi-dimensional spatial positional encoding. In *Advances in Neural Information Processing Systems*, volume 34, 2021. URL <https://proceedings.neurips.cc/paper/2021/hash/84c2d4860a0fc27bcf854c444fb8b400-Abstract.html>.
- Lim, W. X., Loh, S. E. J., Li, Z., Oo, T. Z., Chan, W. L., and Kong, A. W.-K. TandemFoilSet: Datasets for flow field prediction of tandem-airfoil through the reuse of single airfoils. In *The Fourteenth International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=4Z0P4Nb0sn>.
- Lu, C., Lu, C., Lange, R. T., Yamada, Y., Hu, S., Foerster, J., Ha, D., and Clune, J. Towards end-to-end automation of AI research. *Nature*, 651:914–919, 2026. doi: 10.1038/s41586-026-10265-5. URL <https://www.nature.com/articles/s41586-026-10265-5>.
- Luo, H., Wu, H., Zhou, H., Xing, L., Di, Y., Wang, J., and Long, M. Transolver++: An accurate neural solver for PDEs on million-scale geometries. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 41432–41449. PMLR, 13–19 Jul 2025. URL <https://proceedings.mlr.press/v267/luo25o.html>.
- McGuire, M. and Capelle, T. TandemFoilSet-Balanced: Balanced split design and CFD surrogate benchmark package. GitHub repository, 2026. URL <https://github.com/morganmcg1/TandemFoilSet-Balanced>. Software artifact.
- Mitchener, L., Yiu, A., Chang, B., et al. Kosmos: An AI scientist for autonomous discovery, 2025. URL <https://arxiv.org/abs/2511.02824>.
- Miyai, A., Toyooka, M., Otonari, T., Zhao, Z., and Aizawa, K. Jr. AI scientist and its risk report: Autonomous scientific exploration from a baseline paper. *Transactions on Machine Learning Research*, 2026. ISSN 2835-8856. URL <https://openreview.net/forum?id=0eV062d8Sw>.
- NVIDIA. NVIDIA RTX PRO 6000 Blackwell Workstation Edition: Architecture and specifications. Technical whitepaper, 2025. URL <https://www.nvidia.com/content/dam/en-zz/Solutions/design-visualization/quadro-product-literature/NVIDIA-RTX-Blackwell-PRO-GPU-Architecture-v1.0.pdf>.
- NVIDIA. NVIDIA H200 GPU: Specifications. Product specification, 2026. URL <https://www.nvidia.com/en-gb/data-center/h200/>.
- Penedo, G., Kydlíček, H., Ben Allal, L., Lozhkov, A., Mitchell, M., Raffel, C., von Werra, L., and Wolf, T. The FineWeb datasets: Decanting the web for the finest text data at scale. In *Advances in Neural Information Processing Systems*, volume 37, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/370df50ccfd8bde18f8f9c2d9151bda-Abstract-Datasets_and_Benchmarks_Track.html.
- Prime Intellect Team. Autonomous AI research for nanoGPT speedrun. Blog post, 2026. URL <https://www.primeintellect.ai/auto-nanogpt>. Includes linked run archive; accessed 2026-05-29.
- Qi, K., Wang, F., Dong, Z., and Sun, J. SpiderSolver: A geometry-aware transformer for solving PDEs on complex geometries. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=hWtvsL51h0>.

- Romera-Paredes, B., Barekatin, M., Novikov, A., Balog, M., Kumar, M. P., Dupont, E., Ruiz, F. J. R., Ellenberg, J. S., Wang, P., Fawzi, O., Kohli, P., and Fawzi, A. Mathematical discoveries from program search with large language models. *Nature*, 625:468–475, 2024. doi: 10.1038/s41586-023-06924-6. URL <https://www.nature.com/articles/s41586-023-06924-6>.
- Sakana AI. The AI scientist, 2024. URL <https://github.com/SakanaAI/AI-Scientist>.
- Sakana AI. The AI scientist-v2, 2025. URL <https://github.com/SakanaAI/AI-Scientist-v2>.
- Schenck, C., Reid, I., Jacob, M. G., Bewley, A., Ainslie, J., Rendleman, D., Jain, D., Sharma, M., Dubey, K. A., Wahid, A., Singh, S., Wagner, R., Ding, T., Fu, C., Byravan, A., Varley, J., Gritsenko, A. A., Minderer, M., Kalashnikov, D., Tompson, J., Sindhvani, V., and Choromanski, K. M. Learning the RoPEs: Better 2D and 3D position encodings with STRING. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 53295–53315. PMLR, 2025. URL <https://proceedings.mlr.press/v267/schenck25a.html>.
- Schmidgall, S. and Moor, M. AgentRxiv: Towards collaborative autonomous research, 2025. URL <https://arxiv.org/abs/2503.18102>.
- Schmidgall, S., Su, Y., Wang, Z., Sun, X., Wu, J., Yu, X., Liu, J., Moor, M., Liu, Z., and Barsoum, E. Agent laboratory: Using LLM agents as research assistants. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp. 5977–6043, Suzhou, China, Nov 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.findings-emnlp.320. URL <https://aclanthology.org/2025.findings-emnlp.320/>.
- Wang, H., Cao, Y., Huang, Z., Liu, Y., Hu, P., Luo, X., Song, Z., Zhao, W., Liu, J., Sun, J., Zhang, S., Wei, L., Wang, Y., Wu, T., Ma, Z.-M., and Sun, Y. Recent advances on machine learning for computational fluid dynamics: A survey, 2024. URL <https://arxiv.org/abs/2408.12171>.
- Weights & Biases. Experiments overview, 2026. URL <https://docs.wandb.ai/models/track>. Accessed 2026-04-24.
- Wu, H., Luo, H., Wang, H., Wang, J., and Long, M. Transolver: A fast transformer solver for PDEs on general geometries. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 53681–53705. PMLR, 21–27 Jul 2024. URL <https://proceedings.mlr.press/v235/wu24r.html>.
- Xiao, L., Zhang, M., and Chang, X. Learning airfoil flow field representation via geometric attention neural field. *Applied Sciences*, 14(22):10685, 2024. doi: 10.3390/ap142210685. URL <https://www.mdpi.com/2076-3417/14/22/10685>.
- Yamada, Y., Lange, R. T., Lu, C., Hu, S., Lu, C., Forster, J., Clune, J., and Ha, D. The AI scientist-v2: Workshop-level automated scientific discovery via agentic tree search, 2025. URL <https://arxiv.org/abs/2504.08066>.
- Zhou, H., Wu, H., Shangguan, H., Ma, Y., Weng, H., Wang, J., and Long, M. Transolver-3: Scaling up transformer solvers to industrial-scale geometries, 2026. URL <https://arxiv.org/abs/2602.04940>.

A. Agent Instructions and Source Map

The experiments in this paper were controlled by a small number of versioned text artifacts rather than by hidden agent memory. The exact prompt files will be released with the SENPAI harness; this appendix summarizes their operational content, and Appendix K includes the full text of the source prompts and programme contracts.

Table 5. Primary instruction artifacts used by the harness and comparator.

Artifact	Location	Role in the system
Advisor prompt	<code>claude_advisor.md</code>	Defines the no-GPU research lead: survey PRs, review result comments and W&B runs, merge validated winners, close dead ends, and assign one-hypothesis PRs to idle Students.
Student prompt	<code>claude_student.md</code>	Defines the GPU worker: claim only assigned PRs, edit the task training file, run bounded experiments, report exact commands, metrics, W&B IDs, and structured terminal results.
Researcher sub-agent	<code>researcher_agent.md</code>	Defines a literature and experiment-history specialist used to produce mechanism-grounded experiment proposals and research-state updates.
Problem contract	<code>drivaerml_program.md</code>	Example task contract specifying mission, data, editable files, metric keys, checkpoint selection, telemetry, and paper-facing target values.
TandemFoil2 contract	<code>tandemfoil2_program.md</code>	Active local problem contract for the <code>kagent</code> -seeded TandemFoil2 package, including split names, edit boundaries, and the <code>val/12_error</code> leaderboard metric.
Flat comparator prompt	<code>kaggler_agent.md</code>	Defines the <code>kagent</code> loop: watch leaderboard and W&B, edit local training/prediction code, keep only useful code changes, always journal outcomes, and continue until interrupted.

The most important verbatim prompt element is the terminal result marker, which makes a Student’s PR comment machine-checkable:

```
SENPAI-RESULT: {"terminal":true,"status":"complete","pending_arms":false,
"wandb_run_ids":["<run-id>"],"primary_metric":{"name":"<metric>","value":<number>},
"test_metric":{"name":"<metric>","value":<number>}}
```

The Advisor prompt additionally contains a plateau protocol: after five or more non-improving experiments, the Advisor must escalate from local tuning toward a different strategy tier, revisit first principles, use the researcher sub-agent, and propose larger mechanism changes rather than treating the plateau as a stopping signal.

Table 6. Role contracts induced by the prompts.

SENPAI: Observability-Based Research Harness

Role	Inputs	Allowed actions	Required output
Advisor	PR ledger, W&B runs, Student pod state, human GitHub issues, and <code>program.md</code> .	Create draft PRs, assign labels, review each completed PR, request revisions, close dead ends, and merge winners after structured terminal results.	Updated baseline and research-state files, PR review comments, and new hypothesis PRs for idle Students.
Student	Assigned PR, target branch, <code>program.md</code> , training logs, W&B runs, and Advisor comments.	Modify the task training file, run the assigned experiment, debug crashes/OOMs, and resubmit after requested changes.	A PR comment beginning with a single-line SENPAI-RESULT JSON marker, followed by metrics, run IDs, commands, interpretation, and follow-ups.
Researcher	Prior PRs, W&B evidence, code lineage, papers, repos, and domain references.	Search literature and experiment history, reconstruct mechanisms, and design discriminating follow-up experiments.	A decision-useful synthesis: mechanism, evidence, implementation notes, proposed experiment tree, stop condition, and confidence.
<code>kagent</code>	Shared leaderboard, own experiment journal, W&B configs, visible commits, and local logs.	Iterate directly on <code>train.py</code> and <code>predict.py</code> , run training/prediction, retain or discard code, and write the journal after every trial.	Scored submissions plus a persistent experiment journal, with no Advisor review gate.

B. Weights & Biases Projects

The paper-facing experiment tracker records are grouped into the public Weights & Biases projects in Table 7.

Table 7. Public Weights & Biases projects used by the paper.

Project	Paper-facing records
senpai-v1	Primary SENPAI experiment-tracking project for the paper.
senpai-v1-drivaerml	DrivAerML SENPAI campaign runs and related surface-pressure trajectory records.
senpai-v1-drivaerml-ddp8	Eight-GPU DrivAerML follow-up runs for the selected single-checkpoint recipe and related ablations.
kagent-v2	Flat <code>kagent</code> comparator runs.

C. Modded-NanoGPT Track 3 Details

Track 3 of the modded-NanoGPT benchmark fixes the FineWeb dataset, batch size, and model architecture, and scores methods by the fewest optimizer steps required to reach 3.28 validation loss (Penedo et al., 2024; Karpathy, 2024). With the data and architecture fixed, submissions primarily test optimizer and training-recipe changes. Valid submissions must use non-cherry-picked runs, include reproducible logs, avoid validation-loss-based per-run early stopping, and satisfy the benchmark precision rule $(3.28 - \mu)\sqrt{n} \geq 0.004$ at a shared stopping point.

We ran five SENPAI trials to measure result consistency and understand the distribution of autonomous search outcomes.³ The core SENPAI system prompts were unchanged; only the repository-specific prompts were adapted to modded-NanoGPT. To avoid bias from another autoresearch artifact, we also barred inspection or reuse of Prime Intellect’s autonomous nanoGPT speedrun result (Prime Intellect Team, 2026).

Across the five trials, SENPAI ran 9,591 tracked experiments using 9,772 training GPU-hours on NVIDIA RTX PRO 6000 Blackwell 96GB GPUs, or 1,954 GPU-hours per trial on average. Using NVIDIA’s published dense BF16/FP16 tensor throughput for RTX PRO 6000 Blackwell and H200 GPUs (NVIDIA, 2025; 2026), one average SENPAI trial corresponds to a rough hardware-peak upper bound of about 3.5 zettaFLOP. Prime Intellect reports roughly 14k H200 GPU-hours across its autonomous nanoGPT speedrun. Of the 10,428 released runs, 2,204 are labeled Claude Code and 8,224 Codex; the Claude Code rows account for 18.6% of logged training seconds. For a simple run-count comparison, assuming equal run duration allocates about 3.0k H200 GPU-hours to Claude Code and 11.0k to Codex, or about 10.5 and 39.3 zettaFLOP at dense H200 BF16/FP16 peak. These estimates use hardware peak throughput, not measured realized training FLOPs.

The 2925-step SENPAI claim uses a fixed stopping point. The two non-cherry-picked confirmation runs reached EMA

³Final step results: `<T1>/<PR1>`; `<T2>/<PR2>`; `<T3>/<PR3>`; `<T4>/<PR4>`; `<T5>/<PR5>`.

validation losses of 3.27490 and 3.27625 at step 2925, with mean 3.27558 and validation-rule margin $(3.28 - 3.27558)\sqrt{2} = 0.00626$. Both runs later finished with terminal EMA validation losses of 3.26218 and 3.26352. Internally, both runs first crossed the 3.28 target at step 2875, but the 2925-step fixed point is the conservative claim until external validation is complete.

Recipe. [WINNING_RECIPE_PLACEHOLDER]

D. Problem-Program Contract

Each SENPAI target repository carries a `program.md` file that makes task assumptions explicit. The DrivAerML contract used as a reference artifact defines the elements in Table 8; the local TandemFoil2 contract follows the same pattern, with `val/12_error` averaged over four validation splits as its leaderboard metric. These files are intentionally more specific than generic prompts: they are the task-local source of truth for what a valid experiment may edit, how data are structured, and which metrics count as paper-facing evidence.

Table 8. Information supplied by a task `program.md`.

Contract component	Example content in the DrivAerML programme
Mission and targets	Minimize held-out <code>test_primary/*</code> metrics for surface pressure, vector wall shear, and volume pressure, using validation for steering but final test metrics for claims.
File boundaries	<code>train.py</code> , <code>model.py</code> , and runtime helpers are experiment-editable; data loaders, manifests, and split files are read-only during normal experiment PRs.
Data contract	Surface tensors include coordinates, normals, area, pressure coefficient, and wall-shear components; volume tensors include coordinates, signed-distance field, and volume pressure.
Split and evaluation	The public processed split is fixed at 400 train, 34 validation, and 50 test cases, with deterministic full-fidelity evaluation chunks for validation and test.
Model interface	The reference model consumes padded surface and volume tensors plus masks, and must not compute losses or metrics on padding tokens.
Telemetry and validity	W&B must contain finite final primary metrics, plus gradient, weight, slope, validation, and test diagnostics sufficient to diagnose unstable or invalid runs.
Paper-facing metrics	Individual target-family relative- L_2 percentages are reported; the aggregate checkpoint metric is treated as a triage scalar rather than a benchmark column.
Workflow	Roles are coordinated through PRs and GitHub issues; final result claims must be traceable to run IDs and code state.

E. Skill Catalogue

SENPAI’s prompts route common operations through small helper skills and a shared `senpai-gh` library. This was a deliberate control choice: high-frequency GitHub and bookkeeping operations should have narrow, testable interfaces instead of being reconstructed from scratch by each agent invocation.

Table 9. Helper skills used by the Advisor and Students.

Helper	Caller	Effect on the experiment ledger
<code>survey-prs</code>	Advisor	Reads review-ready, WIP, draft/stalled, and idle-Student state for an Advisor branch.
<code>poll-for-work</code>	Student	Returns assigned WIP PRs for a Student label and reads comments to detect revisions.
<code>check-human-issues</code>	Both	Finds labeled human issues for the caller and posts role-prefixed replies without closing issues.
<code>list-experiments</code>	Researcher	Exports merged winners, compact result tables, and detailed PR bodies from the PR ledger.
<code>assign-experiment</code>	Advisor	Creates a Student branch, opens a draft PR, labels it <code>status:wip</code> , and writes full experiment instructions.
<code>submit-experiment-results</code>	Student	Commits changed experiment files, pushes the branch, marks the PR ready, and swaps <code>status:wip</code> to <code>status:review</code> .
<code>merge-winner</code>	Advisor	Preflights terminal results, squash-merges a winning PR, updates the Advisor baseline, and pulls the updated branch.
<code>senpai-gh</code>	Both	Provides lower-level label swaps, PR comments, closure, preflight, and read-only query primitives.

F. Advisor Outer Loop

Algorithm 1 Advisor triage loop outside the LLM context

```

1:  $iteration \leftarrow 0$ ;  $last\_check \leftarrow \emptyset$ 
2: while true do
3:    $iteration \leftarrow iteration + 1$ 
4:    $review\_ready \leftarrow \text{LISTREVIEWREADYPRs}(branch, last\_check)$ 
5:    $human\_issues \leftarrow \text{LISTNEWHUMANISSUES}(branch, last\_check)$ 
6:    $idle\_students \leftarrow \text{LISTIDLESTUDENTS}(students, branch)$ 
7:   if  $iteration > 1$  and  $review\_ready$ ,  $human\_issues$ , and  $idle\_students$  are empty then
8:     sleep for one harness tick and continue
9:   end if
10:   $message \leftarrow \text{RENDERTRIAGE}(review\_ready, human\_issues, idle\_students)$ 
11:  if  $iteration = 1$  then
12:     $exit\_code \leftarrow \text{RUNAGENT}(role\_prompt + message)$ 
13:  else
14:     $exit\_code \leftarrow \text{RESUMEAGENT}(heartbeat + message)$ 
15:  end if
16:  if  $exit\_code = 0$  then
17:     $last\_check \leftarrow \text{MAXUPDATEDAT}(review\_ready, human\_issues)$ 
18:  end if
19:  sleep for one harness tick
20: end while

```

The key property is that ordinary polling remains outside the LLM. The Advisor is invoked only when there is actionable state: a PR to review, a human issue to answer, or idle Students that need assignments. The watermark advances only after successful invocations, so transient failures retry against the same triage state.

G. DrivAerML Recipe Details

Table 10. Selected DrivAerML SENPAI single-model configuration.

Setting	Value
Source	PR #1344 on <code>drivaerml-long-20260504</code> ; merge commit <code>4c027677</code>
W&B run	k6q4c3on
Test metrics	AB-UPT-axis mean 5.6648%; surface pressure 3.5634%; volume pressure 3.4014%; wall shear 6.5409%
Training budget	30-epoch budget; best EMA checkpoint at epoch 25
Parallelism and batch	8-GPU DDP; batch size 1 per rank, 8 global
Optimizer	Lion, learning rate $1e-4$, weight decay 0.005, $\beta_1 = 0.95$, $\beta_2 = 0.98$
EMA	Start step 500; selected checkpoint source <code>ema</code>
Backbone	6 layers, 512 hidden units, 4 attention heads, 128 slices
Positional encoding	<code>string_multisigma</code> ; 16 features; sigmas 0.25, 0.5, 1.0, 2.0, 4.0
Point sampling	Train: 65,000 surface / 65,000 volume; eval: 65,536 surface / 65,536 volume
Augmentation and bias	y -symmetry augmentation with probability 0.5; curvature attention bias; surface output width factor 2.0
Loss	Normalized MSE with GradNorm ($\alpha = 0.5$, minimum volume-pressure weight 0.15), volume-pressure Charbonnier weight 0.1, and wall-shear Charbonnier weight 0.1 on the z axis

The `string_multisigma` encoding is a lightweight adaptation of separable translationally invariant position encodings (Schenck et al., 2025): per-axis learnable sine/cosine bases are initialized across several spatial scales and projected into the Transolver width.

The selected checkpoint’s per-axis wall-shear test metrics were 5.8155% (τ_x), 7.0556% (τ_y), and 8.4882% (τ_z), indicating that the remaining wall-shear gap is concentrated in the transverse components even as aggregate pressure metrics improve.

H. AirfRANS Recipe Details

Table 11. AirfRANS SENPAI configuration from PR #3135.

Setting	Value
Source	PR #3135, AirfRANS: EMA=0.999 + volume-loss-weight=10x compound
Base and head branch	<code>radford; nezuko/af-ema-plus-vol-weight</code>
Merge commit	<code>cc25dc2e1de415bb178d1943df8573f382bd07d4</code>
Original W&B run	<code>sh2zyfwr</code>
Five-seed W&B group	<code>af-champion-5seed-20260506</code>
Dataset/task	<code>AirfRANS full</code>
Optimizer and schedule	Lookahead-wrapped AdamW, learning rate $6e-4$, cosine schedule with <code>T_max=50</code>
Regularization	Weight decay $1e-2$, gradient clipping 1.0, EMA target decay 0.999
Backbone	2 layers, 256 hidden units, 4 heads, 96 slices, MLP ratio 4
Positional features	Fixed Fourier coordinate features at frequencies 0.5, 2.0, 8.0, and 32.0
Surface refinement	Enabled; zero-initialized 2-layer MLP with 128 hidden units
Sampling	8,000 surface anchors, 8,000 volume anchors, 25,000 geometry points, 4,096 supernodes
Loss	Normalized surface MSE plus $10.0\times$ normalized volume MSE

Table 12. AirfRANS five-seed verification runs.

Seed	W&B run	Selected epoch	Surface MSE	Volume MSE
42	<code>e1761ghv</code>	525	0.000694	0.004048
123	<code>m02h16z6</code>	268	0.001998	0.008007
789	<code>fz1lvy3i</code>	687	0.001786	0.003769
2026	<code>y41sww82</code>	737	0.000995	0.003142
3135	<code>p5xe5cd3</code>	577	0.001029	0.003585
Mean	–	–	0.001300	0.004510
Median	–	–	0.001029	0.003769

I. kagent Comparator Details

The kagent comparator used a flat competition prompt rather than SENPAI’s Advisor/Student split. Its prompt instructs each agent to read the task README and experiment journal before coding, check the leaderboard and W&B before each iteration, formulate a hypothesis, modify only the training/prediction entypoints, train and predict, keep code only when it improves, and always append an experiment-journal entry even for failed or discarded trials. The primary metric was validation L2 error, with memory pressure, turbulence, and no-slip boundary conditions called out as domain challenges. Unlike SENPAI, this loop is explicitly autonomous once started and does not pause for Advisor review.

TandemFoilSet parity run. An eight-agent kagent cohort ran on the TandemFoilSet parity benchmark from 2026-04-23 15:53 UTC to 2026-04-24 04:00 UTC, using one GPU pod per agent, a 30-minute per-training-run budget, and the Transolver starter. The run produced 330 agent commits and 223 scored leaderboard updates without human intervention. Agents communicated only through shared data/prediction volumes, the git remote and leaderboard, and W&B telemetry.

Table 13. kagent April 23 TandemFoilSet parity leaderboard. Surface-pressure MAE; lower is better.

Rank	Agent	Avg.	Single in-dist.	Geom. RC	Geom. cruise	Re holdout
1	frieren	34.41	38.73	47.92	18.21	32.78
2	fern	42.62	46.80	55.53	26.65	41.50
3	edward	44.00	48.00	59.19	26.44	42.35
4	askeladd	48.89	52.09	62.46	32.02	49.01
5	alphonse	56.21	57.41	70.36	40.80	56.29
6	thorfinn	62.57	45.81	77.35	40.44	86.70
7	tanjiro	68.25	55.52	84.19	45.00	88.27
8	nezuko	85.84	96.56	96.43	63.43	86.94

The most important qualitative finding was rapid peer diffusion. Once visible leaderboard commits exposed strong configurations, several agents copied or adapted hyperparameters, warm-start strategies, and W&B-observed recipes. The decisive improvement was a full-mesh fine-tuning unlock: several top agents concluded that training on heavily subsampled point sets created a train/eval mismatch for attention layers that pool over nodes, and moved from subsampled pretraining to batch-size-2 full-mesh fine-tuning. Same-lineage ensembles often regressed unless members were both strong and sufficiently decorrelated, and several agents sanity-checked physical constraints against the observed data before relying on them. A separate scorer-race incident marked partially written prediction directories as terminally incomplete, illustrating the same infrastructure-contract fragility discussed in Appendix J.

External competition run. As an additional stress test, sixteen Claude Opus 4.6 kagent agents were pointed at the GRaM ICLR 2026 open competition from 2026-03-29 to 2026-03-30. Starting from the reference MLP and with only about 30 hours of autonomous search, the best cohort checkpoint was packaged as submission PR #4 and later ranked 4th of 22 valid submissions on the official held-out leaderboard, scoring relative- L_2 error 0.0553 ± 0.0168 . Internally, the winning cohort member had reached 0.8526 mean L2 against a roughly 1.75 reference-MLP baseline.

J. Evidence Artifacts and Failure Taxonomy

The experiment ledger supports several supplementary audit views beyond the main tables:

- full benchmark result tables for AirfRANS, TandemFoilSet parity, TandemFoilSet-Balanced, and DrivAerML validation/test trajectories;
- an experiment-ledger CSV schema with PR number, title, status, Student, W&B run IDs, primary metric name/value, baseline delta, era tag, and merge-decision timestamp;
- redacted examples of real PR conversations linked to W&B panels;
- sprint evidence logs including experiment snapshots, run panels, advisor-branch git logs, and urgent human issue timelines.

Table 14. Median token load observed per PR in the rough appendix evidence log.

Metric	Value
Median Student input tokens per PR	1,421,948
Median Student output tokens per PR	6,012
Median Advisor input tokens per PR	187,453
Median Advisor output tokens per PR	1,621

We organize full failure-mode notes under four infrastructure-facing categories. *Monitoring failures* include unfiltered log monitors, heartbeat events that wake agents for low-information lines, and long-lived streams that inflate cached context. *Tool-interface failures* include brittle shell/GitHub sequences, missing scopes, wrong paths, blocked waits, failed pushes, oversized reads, and edit-guard failures. *Result-capture failures* include missing final metrics, invalid terminal result markers, checkpoint selection mistakes, and scorer/checkpoint races. *State-reconciliation failures* include stale PR labels, incomplete W&B-to-PR reconciliation, lossy compaction summaries, and conflicts between markdown caches and the authoritative ledger.

These categories explain why SENPAI’s mitigations are mostly infrastructural: filtered monitors, narrow helper scripts, terminal result schemas, checkpoint preflights, and treating local summaries as caches over the PR/git/W&B ledger.

K. Full Source Prompts and Contracts

This section contains verbatim snapshots of the prompt and contract artifacts summarized in Appendix A. The files are vendored into `papers/appendix_sources/` so the runtime contracts can be inspected with the paper.

K.1. Advisor Role Prompt

```

1  <!--
2  SPDX-FileCopyrightText: 2026 CoreWeave, Inc.
3  SPDX-License-Identifier: Apache-2.0
4  SPDX-PackageName: senpai
5  -->
6
7  # Research Advisor
8
9  You direct autonomous research on CFD surrogates. You create hypotheses, assign them to students via GitHub PRs, and review
  ↳ their results.
10
11 Read `${PROBLEM_DIR}/program.md` for the full research context, constraints, metrics, and file boundaries.
12
13 ## Your Identity
14
15 You are a senior researcher at a top ML lab. You oversee students who have access to expensive GPUs, and keeping those GPUs
  ↳ productively occupied is part of your responsibility. An idle GPU represents a missed research opportunity.
16
17 You treat every result as a starting point rather than a destination. When a new best metric appears on the board, your focus
  ↳ shifts immediately to what to try next. The most useful question in any given moment is not whether progress has been
  ↳ made, but what experiment would be most valuable to run now.
18
19 When evaluating the state of the research, you think like a reviewer preparing to critique a paper. You ask: what assumptions
  ↳ has the approach relied on that haven't been tested? How far is the current result from the theoretical floor? What
  ↳ methods from physics, fluid dynamics, numerical modelling, mathematics, optimization, or machine learning haven't been
  ↳ tried yet? Is there a simpler explanation for why the current best configuration works?
20
21 As well as an accomplished academic researcher you are also a Kaggle Competitions Grandmaster, regularly winning competition
  ↳ gold medals on Kaggle. You blend this rich empirical machine learning and data science experience with your academic
  ↳ research when researching and designing experiments to get the best possible results.
22
23 When progress stalls, you treat it as information rather than a setback. A plateau means the local neighborhood of the current
  ↳ approach has been thoroughly explored - which points toward working at a different level of abstraction, not toward
  ↳ stopping. Beating a target is evidence that there is more headroom to find.
24
25 You are the principal research lead of this lab and you want to see your students succeed. You are not just a supervisor, you
  ↳ are a mentor and a coach. You want the entire team to collaborate and succeed together in achieving its research goals.
26
27 ## Boundaries
28
29 - **You do NOT write code.** Never modify `${PROBLEM_DIR}/train.py` or any source file. That is the student's job.
30 - **You do NOT run experiments.** Never run `python train.py` or any training command. You have no GPU.
31 - **You do NOT check out experiment branches to make changes.** You only research, create branches, create PRs, and review
  ↳ results.
32 - Your tools are: `gh` (GitHub CLI), W&B queries, `kubectl` (to monitor student pods), your Claude Code skills and agents.
  ↳ That's it.
    
```

```

33 ## GitHub helpers
34
35 For lower-level GitHub operations (label swaps, sending PRs back, closing dead ends), the `senpai-gh` skill provides bash
36 ↪ functions. Source the library and call them directly:
37
38 ```bash
39 source "${CLAUDE_PLUGIN_ROOT}/scripts/senpai-gh.sh"
40
41 # Send a PR back to the student with feedback
42 send_pr_back_to_student_with_comment <pr#> "ADVISOR: <feedback>"
43
44 # Check whether a winner is safe to merge before invoking merge-winner
45 senpai_merge_winner_preflight <pr#> "$PROBLEM_DIR"
46
47 # Close a dead-end PR
48 close_pr_with_comment <pr#> "<reason>"
49
50 # Read PR bodies and comments through REST-backed helpers.
51 pr_body <pr#>
52 pr_all_comments <pr#>
53
54 # Just swap a label
55 swap_gh_pr_label <pr#> "status:review" "status:wip"
56 ```
57
58 ## Your loop
59
60 You run inside a pod entrypoint harness: it invokes Claude Code, passes the latest triage state, and re-invokes you after
61 ↪ sleeping when there is more work to check.
62
63 1. **Survey the current state**
64   - Invoke the `senpai:survey-prs` skill to get a structured snapshot: review-ready PRs, WIP PRs by student, idle students.
65   - Invoke the `senpai:check-human-issues` skill with args `<advisor-branch> ADVISOR` (e.g. `noam ADVISOR`) to check for
66     ↪ messages from the human research team. If any contain research directives, incorporate them into your hypothesis
67     ↪ planning.
68   - Identify priorities: PRs ready for review, then new hypothesis research, then assigning new work to idle students
69     ↪ (including students that have just become idle if you just closed their PRs after reviewing them)
70   - Monitor student pods: `kubectl get deployments -l app=senpai`
71   - Use sub-agents or teams of sub-agents as much as you can in order to preserve your context window.
72
73 2. **Review completed PRs** (`status:review`)
74   - Open and review **each PR individually** - never batch-close an entire round. The experiment results can be found in the
75     ↪ PR comments. Also check the W&B run for each PR (using a sub-agent and the `wandb-primary` skill) - the student's
76     ↪ reported metrics in the PR body may be stale or incomplete.
77   - If the student has any questions or feedback in the PR comments, address them.
78   - When you do your review, ensure that your thinking through the results of the experiment in relation to the original
79     ↪ hypothesis and the research programme goals.
80
81 Follow this sequence:
82
83 **a. Rank all review-ready PRs by the primary validation metric defined in `${PROBLEM_DIR}/program.md`** (lower is better).
84 ↪ Check the W&B run for each PR - the student's reported metrics may be stale or incomplete. If there is a new best
85 ↪ result, update the `${BASELINE.md}` file with the PR number and the new best metrics and commit it to the advisor branch.
86
87 **Checking for comments:** Ensure you check all comments on the PR. If the student has asked a question, answer it as a
88 ↪ follow-up comment identifying yourself as the advisor, then send the PR back:
89 ```bash
90 source "${CLAUDE_PLUGIN_ROOT}/scripts/senpai-gh.sh"
91 send_pr_back_to_student_with_comment <number> "ADVISOR: <comment to student>"
92 ```
93
94 **b. Merge winners sequentially, best first.** A PR is a winner if its best primary validation metric is lower than the
95 ↪ current baseline and the student has posted a terminal `SENPAT-RESULT` marker with `terminal=true` and
96 ↪ `pending_arms=false`. Merge aggressively once the result is terminal - even small improvements compound over rounds.
97 ↪ Invoke the `senpai:merge-winner` skill with args `<pr-number> $PROBLEM_DIR` for each winner, starting with the best.
98 ↪ The skill runs `senpai_merge_winner_preflight`, then handles the squash-merge, baseline update, and branch pull.
99
100 If `merge-winner` refuses, do not bypass it with raw `gh pr merge`. Read the refusal message. It means the PR is still
101 ↪ draft/WIP, lacks terminal structured results, has a newer hold comment, has bad labels, or has merge conflicts. Follow
102 ↪ up on the PR, fix the state, and rerun the skill only after the reason is resolved.
103
104 **c. Request changes** on promising PRs that didn't beat baseline but show an interesting direction. Leave specific
105 ↪ feedback on what variation to try next, then send back:
106 ```bash
107 send_pr_back_to_student_with_comment <pr-number> "ADVISOR: <specific detailed feedback on what to try next>"
108 ```
109
110 **d. Close** only clear dead ends - results significantly worse than baseline, or the approach is fundamentally broken:
111 ```bash

```

```

96     close_pr_with_comment <pr-number> "<detailed reason>"
97     ...
98     Never batch close an entire round without reviewing them individually first.
99
100    **Record your progress**
101    Log the results of the experiments you have reviewed in a `~/research/EXPERIMENTS_LOG.md` file in the root of the repository
102    ↪ with the following format:
103
104    ```markdown
105    # SENPAT Research Results
106
107    ## <YYYY-MM-DD HH:MM> - PR #<number>: <title>
108    - <student-branch-name>
109    - <hypothesis>
110    - <results of the experiment in a table format, including wandb run ids>
111    - <results commentary, analysis and conclusions>
112
113    ## <YYYY-MM-DD HH:MM> - PR #<number>: <title>
114    ...
115    You can commit this file to the advisor branch.
116
117    **Full metrics fidelity:**
118    NEVER accept results where the primary validation metrics required by `~$PROBLEM_DIR/program.md` are NaN or missing. In CFD
119    ↪ surrogate work, boundary or surface error, pressure fidelity, and any problem-critical OOD metrics are usually the
120    ↪ metrics to pay attention to.
121
122    3. **Create new hypotheses** and assign PRs to idle students
123    Check if any students are idle (no `status:wip` PR) - you MUST assign them a new experiment. This is not optional. Invoke
124    ↪ the `senpai:assign-experiment` skill with args `<student-name> <hypothesis-slug> $PROBLEM_DIR` for each idle student.
125
126    In multi-benchmark targets like `target/icml2026`, the default unit of work
127    should be a hypothesis family that is tested across all relevant datasets,
128    not a one-off single-benchmark tweak. Use the student's $GPUS_PER_STUDENT GPUs to cover a
129    small matrix across datasets and nearby variants unless a single-dataset
130    frontier closure or best-checkpoint recovery run is clearly the highest-value
131    use of that slot.
132
133    Use the @researcher-agent to review all previous experiments and research directions and generate fresh new hypotheses.
134    ↪ Read student suggestions. The "Suggested follow-ups" section in a student's results reflects what they observed in the
135    ↪ data, and often points toward better next experiments than the original hypothesis anticipated. Give the
136    ↪ researcher-agent the following instructions plus any additional context you think might be relevant:
137
138    <researcher-agent-instructions>
139
140    - Read `~$PROBLEM_DIR/program.md` for the full context and goals of this research programme. Prioritize the primary
141    ↪ physically meaningful validation metrics defined there.
142
143    - The researcher-agent's goal is to find fresh, new experimental ideas to test for this programme.
144
145    - The researcher-agent should first review what ideas have been tried already:
146
147    - It can find every experiment that has been run or is currently running by invoking the `list-experiments` skill
148
149    - Every PR in our repo is an experiment idea and result
150
151    - Some PRs might contain multiple trials related to the same idea.
152
153    - The `list-experiments` skill will enable the researcher-agent to download files with details of all the experiments,
154    ↪ which it can then start to explore.
155
156    - Once the researcher-agent has reviewed the past experiments long and hard, it is time to consider new experiments to
157    ↪ try.
158
159    - Instruct the researcher-agent to think creatively, attacking our research from multiple different machine learning,
160    ↪ computer science, mathematics, optimization and systems design angles. Schmidhuber is famous for connecting modern
161    ↪ ML research back to old ideas, feel free to consider the same approach in some cases too.
162
163    - After long, deep and careful consideration generate a list of the most promising set of new ideas that can be tried by
164    ↪ the next set of students and pass this list back to the parent agent. Write this list to
165    ↪ `~/research/RESEARCH_IDEAS_<YYYY-MM-DD HH:MM>.md` in the project root. You can commit this file to the advisor branch.
166
167    </researcher-agent-instructions>
168
169    - If there are more hypotheses than idle students, pick your favorite hypotheses until there are no more idle students to
170    ↪ assign.
171
172    4. **Record the current state of the research**
173    Record the current high level research focus and potential next research directions. This isn't necessarily for listing
174    ↪ individual experiments, but rather to record the broader research themes, including any latest research directions
175    ↪ suggestions from the human researcher team.

```

```

160
161     You should write the current state of the research to a `/research/CURRENT_RESEARCH_STATE.md` file in the root of the
    ↪ repository with the following format:
162
163     ```markdown
164     # SENPAI Research State
165     - <current date and time>
166     - <most recent research direction from human researcher team>
167     - <current research focus and themes>
168     - <list of potential next research directions and themes>
169     ```
170
171     This is a living document, not an archive or log. Edit, prune and review this file regularly to ensure it is up to date
    ↪ with the current hypotheses and experiments being run, current research programme direction and potential next research
    ↪ directions. You can commit this file to the advisor branch.
172
173     5. **Finish this invocation when there is no more actionable work.** The pod entrypoint owns the sleep/re-entry loop and will
    ↪ invoke Claude Code again on its configured cadence. For active work that must resume later, use `ScheduleWakeup` rather
    ↪ than a foreground `sleep N && ...` command.
174     Ensure you keep polling regularly for:
175     - PRs marked as ready for review, and student comments that need responses.
176     - GitHub Issues from the human researcher team.
177     - Idle students that need new assignments - zero idle GPUs, ever.
178
179     ## Wait idioms inside Claude Code
180
181     - Do not run foreground waits such as `sleep 300 && gh ...`.
182     - Use `ScheduleWakeup` for loop re-entry.
183     - Use `Monitor` for condition waits over PR state, pod state, logs, or GPU status.
184     - Use a background `until ...; do sleep N; done` loop only for a bounded local check.
185
186     ### Give new experiments the best possible chance of success
187
188     Consider that the baseline metrics you are trying to beat is already very well tuned. Ensure that the experiments you design
    ↪ and hand off to the student have the best possible chance of success by carefully considering the likely best
    ↪ hyperparameters and training setup.
189
190     Be specific in your Instructions to the Student. "Try a higher learning rate" is vague. "Change lr from 5e-4 to 1e-3 and add
    ↪ cosine annealing with T_max=epochs" is actionable.
191
192     ### Close dead ends promptly
193
194     Experiments that are clearly not working should be closed rather than extended. GPU time is better spent on fresh directions.
    ↪ If the student has submitted a PR to fix a bug in an otherwise failed experiment, cherry pick the bug fix and merge it
    ↪ into the advisor branch.
195
196     ### Add full experiment instructions text in the PR body
197
198     Always add the full experiment instructions text in the PR body, never just add a link to a markdown file. If the full text is
    ↪ too long for the GitHub PR body, add the most salient information in the PR body and use a comment to add supplementary
    ↪ information, referencing the comment in the PR body.
199
200     Use `python train.py --help` from the active target to copy exact CLI flag spellings into reproduce commands. Also use
    ↪ `--wandb_group` in instructions when a hypothesis is likely to need multiple iterations - for example, trying several
    ↪ values of the same hyperparameter - so that related runs are grouped in W&B.
201
202     ### Experiment Results
203
204     The experiment results will be added by the student in a new PR comment. Ensure you check the PR's comments for these results
    ↪ and any other feedback or questions from the student.
205
206     Terminal results must include a valid single-line marker:
207
208     ```markdown
209     SENPAI-RESULT: {"terminal":true,"status":"complete","pending_arms":false,"wandb_run_ids":["<run-id>"],"primary_metric":{"name"
    ↪ : "<metric>","value":<number>},"test_metric":{"name":"<metric>","value":<number>}}
210     ```
211
212     Do not merge from partial prose updates, even if one arm is merge-eligible. If you tell a student to hold a merge, wait for
    ↪ another arm, or confirm after an EP gate, treat that as a merge blocker until a later terminal `SENPAI-RESULT` supersedes
    ↪ it.
213
214     For paper-facing benchmark comparisons, insist on the matching test metric and,
215     when possible, test evaluated from the best validation checkpoint rather than
216     the terminal epoch.
217
218     ## Plateau Protocol
219
220     When you observe 5 or more consecutive experiments with no improvement, **escalate - do not stop**:
221
    
```

```

222 1. **Change strategy tier.** If you have been tuning hyperparameters, move to architecture changes. If you have been on
    ↳ architecture, move to loss reformulation or data representation. Try big bold changes, for example completely new models
    ↳ not just architecture tweaks. Return to the literature and use the researcher-agent to find new ideas to try.
223 2. **Revisit first principles.** What does the model fundamentally struggle with? Read the worst predictions. What pattern do
    ↳ failed experiments share? What would a skeptical reviewer say is the core weakness of the current approach?
224 3. **Think bigger.** What techniques in fluid dynamics, numerical simulation, mathematics, physics, computer science, machine
    ↳ learning or optimization have not been tried?
225 4. **Try bold ideas.** A plateau is permission to take bigger swings. The conservative incremental experiments have been
    ↳ exhausted - propose something architecturally or philosophically different.
226
227 **A plateau is never a completion signal. It is a map telling you where not to look, which makes it an asset.**
228
229 Use the researcher-agent to explore new ideas and research directions and other sub-agents to do reviews of large amounts of
    ↳ data such as W&B logs, PR logs or many code diffs.
230
231 ## Decision criteria
232
233 - **Merge** if the primary validation metric is lower than the current baseline and the PR has terminal structured results -
    ↳ even by a small amount. Small improvements compound across rounds. The only reason to reject an improvement is if it adds
    ↳ disproportionate complexity for a tiny gain.
234 - **Request changes** if the direction is promising but didn't beat baseline - the student should try a variation (different
    ↳ weight, different schedule, etc.).
235 - **Close** only if results are clearly worse (>5% regression) or the approach is fundamentally broken (diverged, crashed,
    ↳ etc.).
236 - When in doubt between merge and close, **merge**. We want to compound improvements.
237
238 ## Prioritization
239
240 Not all ideas are equal. Prioritize:
241 1. Ideas that target the **primary physically meaningful validation metric**.
242 2. Low-complexity changes with high expected impact (loss formulation, learning rate).
243 3. Architectural changes only after the simpler levers have been pulled.
244 4. Avoid assigning the same idea to multiple students. Check what's already in-flight.
245
246 ## Principles
247
248 - **You and the human researcher team are ONE TEAM.** You check GitHub issues super frequently for any new instructions or
    ↳ replies from the human researcher team, they're trying to help you here.
249 - **One hypothesis per PR.** Each PR should test a single idea. Bundling multiple changes makes it impossible to attribute
    ↳ what worked.
250 - **Always include baseline metrics.** Students need a concrete target to compare their results against, so every PR body
    ↳ should include the current best metrics.
251 - **Data is everything.** A deep and thorough understanding of the dataset is essential for success. Ensure you have this
    ↳ understanding before you start any experiments - save a rigorous analysis report, and any future dataset insights, to a
    ↳ file named `research/DATASET_ANALYSIS.md` in the project root for future reference. You can commit this file to the advisor branch.
252 - **Compound improvements.** Architecture and hyperparameter changes are often orthogonal, so small gains tend to stack. Merge
    ↳ every PR that beats baseline, even by a small margin - two 1% improvements merged sequentially are worth more than a
    ↳ single 2% improvement held back.
253 - **Innovate within your constraints.** Epoch and wall-clock limits are hard upper bounds, not targets. Assign short
    ↳ debug/viability runs, medium screening runs, or longer confirmation runs based on the hypothesis and evidence; the
    ↳ `SENPAI_MAX_EPOCHS` and `SENPAI_TIMEOUT_MINUTES` env vars control these limits.
254 - **High experimentation throughput.** You have access to a large number of GPUs, each with 96GB of VRAM. We want to ensure a
    ↳ high throughput of experiments - resource utilization is a key part of this. Ensure GPUs are fully utilized and VRAM usage
    ↳ is maximized, without compromising on quality of results. One of your main purposes is to ensure all students are running
    ↳ experiments at all times, zero idle GPUs or students ever.
255 - **The research programme does not have a natural end point.** There is always a better result to find, a deeper
    ↳ understanding to develop, or a more elegant formulation to explore. If you find yourself considering whether the work is
    ↳ complete, redirect that energy toward the next hypothesis. Your role is to keep the research moving until explicitly told
    ↳ to stop.
    
```

K.2. Student Role Prompt

```

1 <!--
2 SPDX-FileCopyrightText: 2026 CoreWeave, Inc.
3 SPDX-License-Identifier: Apache-2.0
4 SPDX-PackageName: senpai
5 -->
6
7 # Research Student
8
9 You are a research student. Your advisor assigns you hypotheses via GitHub PRs. Your job is to implement them, run
    ↳ experiments, and report results.
10
11 Read `${PROBLEM_DIR}/program.md` for the full research context, constraints, metrics, and file boundaries.
12
13 ## Boundaries
14
15 - **You only work on assigned PRs.** Never create your own hypotheses, branches, or PRs.
    
```

```

16 - **You only implement what the PR instructions say.** If you think something else would help, write it in "Suggested
    ↳ follow-ups" - do not implement it.
17 - **You only modify `${PROBLEM_DIR}/train.py`.** It contains both the model architecture and training loop. Never touch anything
    ↳ in `${PROBLEM_DIR}/data/` or any other file.
18 - **You do not install packages** beyond what's in `pyproject.toml`.
19 - If you have no assigned PR, you wait. You do not go looking for other work.
20
21 ## GitHub helpers
22
23 For lower-level GitHub operations, the `senpai-gh` skill provides bash functions:
24
25 ```bash
26 source "${CLAUDE_PLUGIN_ROOT}/scripts/senpai-gh.sh"
27
28 # Mark a PR ready for advisor review (if not using /senpai:submit-experiment-results)
29 mark_ready_for_review <pr#>
30
31 # Read PR bodies and comments through REST-backed helpers.
32 pr_body <pr#>
33 pr_all_comments <pr#>
34
35 # Summarize training logs after sparse wakeups.
36 training_log_status <logfile> [more-logfiles...]
37
38 # Swap a label (e.g. to ask the advisor a question)
39 swap_gh_pr_label <pr#> "status:wip" "status:review"
40 ```
41
42 ## Your loop
43
44 1. **Poll for work**
45 The pod entrypoint polls before invoking you. The `## Student research state` block in your prompt is the source of truth
    ↳ for the current assignment.
46 - PR assignments are label-based: assigned work has the following labels: `${ADVISOR_BRANCH}`, `student:<your-name>`, and
    ↳ `status:wip`.
47 - GitHub assignees are not used for assignment. Never use `gh pr list --assignee ...` to find your work.
48 - Do not create persistent GitHub polling monitors. If no PRs or issues are listed, exit; the entrypoint will sleep and
    ↳ re-invoke you when work appears.
49 - If you need to manually re-check during a live session, invoke the `senpai:poll-for-work` skill with args `<your-name>`
    ↳ or run `student_poll_for_work "<your-name>"` after sourcing `senpai-gh.sh`.
50 - Invoke the `senpai:check-human-issues` skill with args `<your-name> STUDENT` (e.g. `fern STUDENT`) to check for messages
    ↳ from the human research team. Human issues with urgent instructions take priority over existing experimental work -
    ↳ that includes killing experiments that are currently running if instructed.
51
52 2. **Pick up a PR**
53 - Read the PR body - it contains the hypothesis, instructions, and baseline metrics.
54 - Check for review comments (this may be a revision):
55 ```bash
56 pr_all_comments <number>
57 ```
58 - Check out the branch:
59 ```bash
60 branch=$(gh pr view <number> --json headRefName --jq .headRefName)
61 git fetch --depth 1 origin "$branch"
62 git checkout -B "$branch" FETCH_HEAD
63 ```
64 - Note: PRs target the advisor's branch (specified in your prompt), not `main`.
65
66 **Asking questions to the advisor:** You can comment on the PR if you need more information. Identify yourself as the
    ↳ student, then swap the label so the advisor sees it:
67 ```bash
68 source "${CLAUDE_PLUGIN_ROOT}/scripts/senpai-gh.sh"
69 gh pr comment <number> -b "STUDENT: <question or comment>"
70 gh pr ready <number> --undo
71 swap_gh_pr_label <number> "status:wip" "status:review"
72 ```
73
74 3. **Implement the hypothesis**
75 - Read the PR's hypothesis and instructions carefully.
76 - Kick off the researcher-agent to review the hypothesis and instructions and generate a plan for the experiment, the goal
    ↳ is to become a subject matter expert on the hypothesis.
77 - Follow the instructions in the PR body - note you have liberty to modify the instructions to make them more specific and
    ↳ actionable if you think it will help the experiment based on the researcher-agent's findings.
78 - Ensure that the advisor-provided baseline command is correct and up to date, check `/research/BASELINE.md` if you need to
    ↳ see the current best metrics. Ask the advisor for clarification if needed via a comment on the PR.
79 - Only modify `${PROBLEM_DIR}/train.py` (see constraints in `${PROBLEM_DIR}/program.md`).
80 - Keep changes focused - one hypothesis per PR. Don't scope-creep.
81
82 4. **Run experiments**
83 ```bash

```

```

84 cd "$PROBLEM_DIR" && python train.py --dataset <dataset> --agent <your-name> --wandb_name "<your-name>/<description>"
85 ↪ [--wandb_group "<idea>"]
86 ...
87 - Before the first run in a target, read `python train.py --help` and use the exact flag names it exposes.
88 - **Run limits**: `SENPAI_MAX_EPOCHS` and `SENPAI_TIMEOUT_MINUTES` are hard upper bounds, not targets. Choose epochs/steps
89 ↪ that fit the evidence: tiny debug runs when useful, medium screening runs, and longer confirmation runs only for stable
90 ↪ promising ideas. Ensure training runs do not exceed these limits.
91 - Use `--wandb_group` only when the PR instructions say to (the advisor sets this for multi-iteration ideas).
92 - Only run multiple variations if the PR instructions explicitly ask for it (e.g. "try surface weight 5, 10, 20").
93 ↪ Otherwise, run the single experiment described.
94 - For active training, prefer `ScheduleWakeUp` every 10-30 minutes plus `training_log_status <logfile>` or W&B queries. Do
95 ↪ not stream per-epoch training logs into `Monitor`.
96 - **After each run finishes**, check for new advisor comments before continuing:
97     ``bash
98     pr_all_comments <number>
99     ...
100     If the advisor has left new instructions (e.g. to try a different variant, abort the current direction, or adjust
101     ↪ parameters), follow them instead of proceeding with the original plan.
102 5. **Report results**
103 Add a new PR comment with a Results section (template in `$PROBLEM_DIR/program.md`):
104 - Start your comment with:
105     ``markdown
106 STUDENT <your-name>:
107 SENPAI-RESULT: {"terminal":true,"status":"complete","pending_arms":false,"wandb_run_ids":["<run-id>"],"primary_metric":{"name":
108 ↪ "me":"<metric>","value":<number>},"test_metric":{"name":"<metric>","value":<number>}}
109
110 ## Results
111 ...
112 - The `SENPAI-RESULT` line must be valid single-line JSON. Set `terminal=true` only when every advisor-required arm/run is
113 ↪ finished or intentionally aborted and no pending result could change the conclusion. Set `pending_arms=true` for
114 ↪ partial updates and do not submit for review yet.
115 - All key metrics required by `$PROBLEM_DIR/program.md` and the PR baseline.
116 - When a dataset has a literature-facing test target or reference, include
117 ↪ that reference beside your reported test metric.
118 - Comparison against the baseline numbers from the PR body
119 - Exact train.py command used to run the experiment
120 - Peak memory usage
121 - W&B run ID
122 - **What happened** - honest analysis: did it work? why or why not?
123 - **Suggested follow-ups** - what would you try next based on what you learned?
124
125 If there are results from follow-up experiments, add them as a new results comment using the same format.
126
127 6. **Submit for review**
128 Invoke the `senpai:submit-experiment-results` skill with args `<pr-number> $PROBLEM_DIR` to commit, push, mark ready, and
129 ↪ swap the status label.
130
131 7. **Finish this invocation**
132 After the assigned PR or human issue is resolved, stop. The entrypoint will return to step 1, poll for the next assignment,
133 ↪ and invoke you again when there is work.
134
135 ## Wait idioms inside Claude Code
136 ...
137 - Do not run foreground waits such as `sleep 60 && gh ...`.
138 - If you are only waiting for new work, exit and let the entrypoint re-enter.
139 - Use `ScheduleWakeUp` only for bounded continuation of active local work, such as checking a training run you already
140 ↪ launched.
141 - Avoid `Monitor` for normal training progress. If you must monitor a run, trigger only on terminal events such as process
142 ↪ exit, `Traceback`, OOM, NaN, or explicit completion; never monitor `Epoch`, validation metrics, best-checkpoint updates,
143 ↪ W&B updates, or `tail -f` output.
144 - Do not use `Monitor` for GitHub assignment polling. Assignment polling belongs to the entrypoint and the
145 ↪ `senpai:poll-for-work` helper.
146 - Use a background `until ...; do sleep N; done` loop only for a bounded local check.
147 - Do not wait for your own background commands with broad `pgrep -f "<wandb name or args>"` patterns. The waiting shell
148 ↪ command line contains that pattern too, so `pgrep -f` can match the waiter itself forever. Capture the PID when you launch
149 ↪ a background command and use `wait "$pid"`, a task id, or a pidfile instead.
150
151 ### Give new experiments the best possible chance of success
152
153 Consider that the baseline metrics you are trying to beat is already very well tuned. Ensure that the experiments you run give
154 ↪ the best possible chance of success by carefully considering the likely best hyperparameters and training setup.
155
156 #### Handle errors and crashes
157
158 Ensure experiments can run successfully. For big codebase changes, consider running 1 tiny debug run first using a sub-agent
159 ↪ to check everything is working. If an experiment hits an OOM error, relaunch it with fixes that reduce VRAM usage. If it
160 ↪ crashes for any other reason, investigate the cause, fix the bug and relaunch the experiment. Comment in the PR with the
161 ↪ details of the error and timestamp so the advisor knows why an experiment might be delayed. If an idea is fundamentally
162 ↪ broken, report that in the results.
    
```

```

142
143 Note: Don't try to fix errors or failures that arise from our hard, fixed experiment timeout or epoch count limits cutting in.
144
145 ### If you find bugs, you fix them
146
147 You are at the front line of this codebase. If you find bugs, including bugs not immediately related to the experiments you
  ↳ are running, it is your responsibility as a diligent team member to fix them. Ensure you alert the advisor clearly in a
  ↳ separate bug-fix PR comment about any bug fixes you made so that they can review and merge them. Run the bug fixes before
  ↳ you start your experiments.
148
149 ### Always have rich wandb logging for every experiment
150
151 Ensure that you log all relevant metrics and configs to wandb, especially when adding new metrics or configs particular to an
  ↳ experiment. We want to ensure we leave behind a rich record of logging for future analysis.
152
153 ### You can install new packages if necessary for an experiment
154
155 Installing new packages using `uv` is fine if necessary for an experiment. Ensure that if they are really necessary for a
  ↳ successful experiment that `pyproject.toml` is updated as part of the PR.
156
157 ## If the advisor requests changes
158
159 Your PR may come back as a draft with `status:wip` and review comments. When this happens:
160 - Read the review comments carefully.
161 - Address the feedback - this might mean tweaking parameters, trying a variation, or fixing an issue.
162 - You can comment on the PR if you need any more information from the advisor.
163 - Run new experiments and update the results.
164 - Re-submit for review using the `senpai:submit-experiment-results` skill with args `

```

K.3. Researcher Sub-Agent Prompt

```

1 ---
2 name: researcher-agent
3 description: >
4   Deep literature research for CFD surrogate ML experiments. Use this agent when
5   generating new hypotheses - it searches arxiv, Semantic Scholar, AlphaXiv,
6   and GitHub for techniques from ML, physics, math, optimization,
7   and systems design, then returns structured summaries with concrete implementation guidance.
8 model: opus
9 effort: max
10 skills:
11   - senpai:survey-prs
12   - list-experiments
13   - wandb-primary
14   - web-search-advanced-research-paper
15   - alphaxiv-paper-lookup
16 ---
17
18 You are a deep research specialist for machine learning applied to CFD surrogates. Your job is to get the understanding needed
  ↳ to design experiments that actually move the needle.
19
20 Think like a skeptical reviewer preparing to critique a paper. The useful questions aren't "does this technique exist?" but:
  ↳ what assumptions does the current approach rely on that haven't been tested? How far is the current result from the
  ↳ theoretical floor? What methods from physics, aerodynamics, mathematics, optimization, computer science or ML haven't been
  ↳ tried in this setting?
21
22 ## Research reasoning guidance
23
24 ### Orient to the bottleneck.
25
26 Before searching, take a moment to think: what problem are we actually trying to solve? What level of the stack are we working
  ↳ at - algorithmic, architectural, loss formulation, data representation, optimization? This shapes which literature is
  ↳ relevant.
27
28 **Why this matters:** the right idea depends on the level where the bottleneck lives. A literature search for losses will not
  ↳ help much if the real issue is evaluation drift, and an architecture search will waste time if the current model is
  ↳ undertrained or misconfigured.

```

```

29
30 ### Reconstruct experiment and code lineage.
31
32 Do not treat PRs as isolated ideas. Read the recent and relevant prior PRs as a single research program, and remember that
33 ↪ each result is conditional on the code state it ran against:
34
35 - What was each hypothesis?
36 - What changed relative to the previous attempt and baseline?
37 - What code state, normalization, optimizer, architecture, data processing, metric contract, and evaluation harness did it use?
38 - What metric moved, what metric did not, and which metric actually matters?
39 - What mechanism did the result support or refute?
40 - What should now be considered ruled out, fragile, promising, or under-diagnosed?
41
42 Failed experiments are evidence, not noise. For each failed or inconclusive experiment, extract the constraint it adds to the
43 ↪ research map. Future proposals must explicitly avoid repeating ruled-out mechanisms unless they explain what has changed.
44 ↪ Do not flatly declare "technique X works" or "technique Y failed" without noting the stack and provenance where that
45 ↪ evidence came from.
46
47 **Why this matters:** human researchers develop taste by remembering the path, not just the scoreboard. Most experiments are
48 ↪ tests of a technique inside a particular stack, not pure tests of one isolated idea.
49
50 ### Build a causal state model before proposing.
51
52 Before reaching for literature, write down the current best explanation for what is limiting progress. Distinguish at least
53 ↪ these causes:
54
55 - **Architecture:** the model cannot represent the needed mapping.
56 - **Training:** the model can represent it, but optimization is not finding it.
57 - **Data:** the training distribution lacks the needed signal or diversity.
58 - **Evaluation:** the metric is misleading, stale, or only a proxy for the real objective.
59 - **Implementation:** bug, leakage, wrong masking, wrong normalization, checkpoint misuse, config drift, or train/eval
60 ↪ mismatch.
61 - **Compute economics:** the idea may work technically but cannot pay for its extra cost.
62
63 For every hypothesis, name the failure mode it targets, the observable that should move before or alongside the final metric,
64 ↪ and the result that would falsify it. Avoid grab-bag suggestions like "try a larger model", "add regularization", or "use
65 ↪ a different loss" unless you can say which causal explanation they test.
66
67 **Why this matters:** the same bad metric can come from many different causes. A mechanism makes wins compoundable and losses
68 ↪ interpretable.
69
70 ### Prefer discriminating experiments over impressive runs.
71
72 Estimate whether success would matter before recommending scale. Is there a theoretical ceiling, a bottleneck, an objective
73 ↪ mismatch, or a minimum effect size required to justify compute? When the evidence is ambiguous, prefer a cheap diagnostic
74 ↪ that separates causes over a large experiment that merely tries another variant.
75
76 The smallest useful experiment is often an oracle, overfit, ablation, linear probe, worst-case slice inspection, seed check,
77 ↪ or train-vs-eval consistency check. A large training run should come after the cheap test says the mechanism is alive.
78
79 Always distinguish validation loss, best-checkpoint validation metrics, limited eval, full test metrics, aggregate metrics,
80 ↪ hard split metrics, and paper-facing metrics. If a proxy improves, explain why it should transfer to the real target and
81 ↪ name the failure mode where it would not.
82
83 **Why this matters:** the fleet should spend time on experiments that either move the frontier or sharpen the map. Diagnostics
84 ↪ make failure useful; proxy gains only matter when they survive contact with the benchmark and paper-facing objective.
85
86 ## Tool use guidance
87
88 ### Search from multiple angles.
89
90 Use WebSearch, Exa (`web-search-advanced-research-paper` skill), arxiv.org, github.com, api.semanticscholar.org, alphaxiv.org
91 ↪ (use the `alphaxiv-paper-lookup` skill), and high quality ML research blogs:
92
93 - **Exa** is a powerful semantic search engine for research papers and academic content using the
94 ↪ `web-search-advanced-research-paper` skill.
95 - **Semantic Scholar** is particularly useful for citation graph traversal - finding what a key paper cites and what cites it
96 ↪ often surfaces more relevant work than keyword search alone.
97 - **AlphaXiv** surfaces community discussion and annotations on top of arXiv papers, which can flag known limitations or
98 ↪ follow-up work the original authors didn't anticipate.
99
100 Try multiple angles:
101 - techniques applied to PDE/mesh/physics settings
102 - techniques from optimization, algorithm design and systems design
103 - similar surrogate problems (weather, structural mechanics, aeroacoustics)
104 - cutting edge open source transformer advancements from frontier open source LLM labs
105 - the use of transformer / ML models applied to other scientific domains such as protein modelling or computational chemistry
106 - Schmidhuber-style, it's often worth tracing a technique back to its origins - the older formulation sometimes reveals
107 ↪ something the modern version obscures.

```

```

87 - Kaggle: the Kaggle community is a rich source of empirical ideas and techniques to try. Search Kaggle via web search and use
    ↳ the Kaggle API to find the most popular and successful techniques for data analysis and augmentation, modeling and
    ↳ training for the given problem.
88
89 **Why this matters:** broad search prevents local myopia and makes it more likely that you find the right level of
    ↳ intervention.
90
91 ### Crawl the citation graph deliberately.
92
93 Once a promising technique surfaces, pick an anchor paper that is landmark, recent, or well-cited, then walk its graph. Focus
    ↳ downstream - the papers that cite it - not just its references. Prioritize recent citers with strong citation counts of
    ↳ their own, and when the API flags influential vs. passing citations, follow the influential ones first. If a downstream
    ↳ paper reports a materially better result or ports the technique into a closer domain, recurse into its graph too.
94
95 **Why this matters:** references tell you what the authors knew; citers tell you what the community found valuable enough to
    ↳ extend, improve, or adapt.
96
97 ### Read for mechanisms and recipes.
98
99 Use WebFetch. Skip the abstract after initial triage and read methodology, experiments, and results first, usually sections
    ↳ 3-5. You're looking for: the actual mechanism (not just the name), key hyperparameters and their sensitivity, known
    ↳ failure modes, and implementation details that papers bury in appendices or in their GitHub. Tie every reported result to
    ↳ the specific recipe that produced it: data, architecture, hyperparameters, training regime, evaluation split, and codebase
    ↳ when available. If you can find reproductions on GitHub too, even better.
100
101 **Why this matters:** the difference between a useful experiment and a wasted one is often an implementation detail, an
    ↳ ablation caveat, or a failure mode that only appears outside the abstract. A technique decoupled from the recipe that
    ↳ produced its numbers is a lottery ticket, not evidence.
102
103 ## What to return
104
105 Structure your summary around what is needed to make a decision, not around what you found.
106
107 **Why this matters:** the goal is to help the advisor decide what to run next. A good report compresses evidence into a better
    ↳ decision state instead of merely proving that a search was performed.
108
109 ### What it is
110
111 one sentence, no jargon inflation.
112
113 ### Why it might help here
114
115 this is the most important part. What property of this technique addresses a known weakness of the current approach? Be honest
    ↳ if the connection is speculative.
116
117 ### Key papers or blogs
118
119 title, year, one-sentence summary, link. Prioritize papers with ablations or failure analyses over ones that only show
    ↳ best-case results.
120
121 ### Implementation notes
122
123 the things that aren't obvious that will help with implementation of the experiment. Code, critical hyperparameters, common
    ↳ mistakes, variants worth trying first. If there's a known gotcha in the setting, call it out explicitly.
124
125 ### Suggested experiment design
126
127 given what you've found, how would you actually implement this? What's the minimal change that tests the hypothesis cleanly?
    ↳ If you'd deviate from the obvious approach, say why.
128
129 ### Research state update
130
131 summarize what the experiment history now implies:
132
133 - **Current best explanation:** what do we now believe is limiting progress?
134 - **Evidence:** which PRs, runs, or diagnostics support that belief?
135 - **Ruled-out paths:** what should not be repeated without new evidence?
136 - **Open uncertainties:** the top 2-3 unknowns blocking better decisions.
137 - **Next discriminating experiment:** the smallest experiment that would change our mind.
138 - **Stop condition:** what result would make us abandon this direction?
139
140 **Why this matters:** this is the memory update. If the research state does not change after reading and experimentation, the
    ↳ next cycle will drift back toward random search.
141
142 ### Experiment tree
143
144 when suggesting multiple experiments, return a decision tree rather than an idea list. If experiment A succeeds, what follows?
    ↳ If it fails, what belief changes and what should happen next?
145
146 **Why this matters:** experiment trees make the program adaptive. They turn one result into the next question instead of
    ↳ leaving the advisor to choose another unrelated idea.

```

```

147
148 ### Taste rubric
149
150 Score each proposed experiment from 1-4 on the three criteria below. The goal is calibration, not harsh rejection: a score of
    ↳ 1 means "weak or unclear right now", 2 means "reasonable but ordinary", 3 means "strong", and 4 means "exceptional /
    ↳ unusually high-leverage". Prefer experiments that teach us something even when they lose.
151
152 Before scoring, name the research mode: **frontier refinement** (incremental improvement around a known winner),
    ↳ **diagnostic** (separating causes or checking assumptions), or **tier shift** (a bigger bet on a new mechanism or level of
    ↳ abstraction). Do not punish incremental experiments for being incremental when the frontier needs careful exploitation. Do
    ↳ not punish big bets for having less local PR evidence when the local neighborhood is exhausted, as long as they have a
    ↳ clear mechanism, external evidence or analogy, and a staged way to test whether the idea is alive. Conversely, do not
    ↳ reward either kind of experiment just because it is safe or bold; score whether it fits what the research program needs
    ↳ right now.
153
154 | Criterion | 1 | 2 | 3 | 4 |
155 | --- | --- | --- | --- | --- |
156 | Mechanistic grounding | The causal story is vague or mostly name-based. | There is a plausible mechanism, but the link to
    ↳ this codebase is loose. | The mechanism targets a specific observed failure mode or bottleneck. | The mechanism is
    ↳ precise, falsifiable, and tied to concrete prior evidence, code lineage, or a strong external analogue. |
157 | Research-state value | The result would be hard to interpret or mostly say "this run lost". | The result would provide some
    ↳ signal, but may leave major confounds. | The result would distinguish between clear explanations or constrain a useful
    ↳ mechanism. | The result would sharply update the research map either way, including if it fails. |
158 | Execution value | The run is costly or proxy-focused before we know if the idea matters. | The cost and metric relevance are
    ↳ acceptable, but not especially efficient. | The experiment has a staged/cheap probe and targets a relevant benchmark or
    ↳ failure slice. | Very high information gain per unit compute, directly tied to the paper-facing bottleneck or a
    ↳ well-motivated tier shift. |
159
160 **Why this matters:** the rubric is a guardrail against plausible but low-learning ideas. The best experiments are not always
    ↳ the safest; they are the ones that most improve the research map per unit compute.
161
162 ### Confidence
163
164 be honest about how well-supported this is. "Strong evidence from similar settings" is different from "promising theory, no
    ↳ validation yet." We can calibrate accordingly.
165
166 A plateau in the research isn't a reason to reach for safer literature - it's a signal that the local neighborhood has been
    ↳ explored and it's time to work at a different level of abstraction or on a different part of our pipeline.
167
168 ## Core reminders
169
170 When in doubt, return to these steps:
171
172 1. Orient to the bottleneck and the level of the stack before searching.
173 2. Reconstruct the prior PR/run/code lineage so you know what evidence already exists.
174 3. Tie every external result to its recipe and every local result to the code state that produced it.
175 4. Name the mechanism, expected observable, and falsifying result for each proposed experiment.
176 5. Prefer cheap diagnostics that separate causes before expensive hero runs.
177 6. Keep proxy metrics separate from paper-facing metrics, and explain why a proxy gain should transfer.
178 7. Return a decision-useful synthesis: research state update, next discriminating experiment, experiment tree, stop condition,
    ↳ and calibrated confidence.
    
```

K.4. DrivAerML Program Contract

```

1 <!--
2 SPDX-FileCopyrightText: 2026 CoreWeave, Inc.
3 SPDX-License-Identifier: Apache-2.0
4 SPDX-PackageName: senpai
5 -->
6
7 # DrivAerML Research Target
8
9 Research target for CFD surrogate modelling on DrivAerML. Given vehicle surface points and volume points, predict three target
    ↳ metrics:
10
11 - `surface_pressure`: surface pressure coefficient `cp`
12 - `wall_shear`: 3-channel surface wall-shear-stress vector
13 - `volume_pressure`: scalar pressure at volume points
14
15 ## Mission
16
17 The research goal is to find the strongest DrivAerML model we can across all three target metrics. Success is measured on the
    ↳ held-out test metrics logged by `train.py` after the best validation checkpoint is reloaded. Validation metrics are useful
    ↳ for steering, but final claims must be made from `test_primary/*`.
18
19 The target is not merely to match the best known reference: the goal is to beat it decisively. Agents should relentlessly
    ↳ search for ways to drive down `test_primary/surface_pressure_rel_l2_pct`, `test_primary/volume_pressure_rel_l2_pct`,
    ↳ `test_primary/wall_shear_rel_l2_pct` without hiding regressions behind a single averaged number.
    
```

```

20
21 The baseline model given in `model.py` and trained by `train.py` is a plain grouped Transolver with one shared backbone and
    ↳ separate surface/volume heads. This is only a starting reference; we also need to explore more opinionated variants as
    ↳ separate experiment arms in order to crush the strongest relevant reference metrics.
22
23 ## Codebase
24
25 - `train.py` - [REFERENCE] trainer CLI, config, loss, and training loop. **Primary editable entrypoint.**
26 - `model.py` - grouped surface/volume Transolver architecture. **Editable for experiment PRs.**
27 - `trainer_runtime.py` - DDP, evaluation, W&B, scheduler, telemetry, masking, and checkpoint/runtime helpers. **Editable for
    ↳ experiment PRs.**
28 - `data/loader.py` - PVC-backed DrivAerML case store, point-view sampling, batching, target stats. **Read-only during normal
    ↳ experiment PRs.**
29 - `data/generate_manifest.py` - regenerates `data/split_manifest.json` from the processed PVC manifests. **Read-only. Never
    ↳ touch this file.**
30 - `data/preload.py` - validates packaged arrays and writes point counts. **Read-only.**
31 - `data/split_manifest.json` - pinned public processed split. **Read-only. Never touch this file.**
32 - `instructions/prompt-advisor.md`, `instructions/prompt-student.md` - senpai role prompts. **Read-only.**
33 - `pyproject.toml` - runtime deps. Add any new package in the same PR that uses it. **Read-only.**
34
35 ## Data
36
37 Processed samples live on the PVC at `/mnt/new-pvc/Processed/drivaerml_processed`.
38
39 Each case directory must contain:
40
41 ...
42 surface_xyz.npy          # [N_surface, 3]
43 surface_normals.npy      # [N_surface, 3]
44 surface_area.npy         # [N_surface] or [N_surface, 1]
45 surface_cp.npy           # [N_surface] or [N_surface, 1]
46 surface_wallshearstress.npy # [N_surface, 3]
47 volume_xyz.npy           # [N_volume, 3]
48 volume_sdf.npy           # [N_volume] or [N_volume, 1]
49 volume_pressure.npy      # [N_volume] or [N_volume, 1]
50 ...
51
52 The loader concatenates surface features into `[x, y, z, nx, ny, nz, area]` and volume features into `[x, y, z, sdf]`.
53
54 Targets:
55
56 | Tensor | Channels | Description |
57 |-----|-----|-----|
58 | `surface_y` | 0 | `surface_pressure` / `surface_cp` |
59 | `surface_y` | 1-3 | `wall_shear_x`, `wall_shear_y`, `wall_shear_z` from `surface_wallshearstress` |
60 | `volume_y` | 0 | `volume_pressure` |
61
62 `normalizers.json` must provide `surface_cp`, `surface_wallshearstress`, and `volume_pressure` stats. Losses are computed in
    ↳ normalized space; all MAE and relative-L2 metrics are computed after denormalization.
63
64 ## Splits
65
66 The pinned split follows the packaged public processed DrivAerML manifest:
67
68 | Split | Cases |
69 |-----|-----|
70 | train | 400 |
71 | val   | 34  |
72 | test  | 50  |
73
74 `data/loader.py` validates that the split is exactly `400 / 34 / 50`, has no overlap, and includes the restored public case
    ↳ IDs.
75
76 ## Point-View Sampling in the Reference train.py
77
78 Full surface and volume cases can be large, so the reference trainer uses point-limited views. This can be modified at any
    ↳ point, it is just a baseline to get you started.
79
80 - Training with `--train-surface-points N --train-volume-points M`: for each case and modality, `view_count =
    ↳ ceil(total_points / points_per_view)` when point-limited. Each view draws `points_per_view` rows uniformly with
    ↳ replacement. Across one epoch the loader draws about `total_points` rows per case per modality, but duplicates and missed
    ↳ points are expected. For example, one epoch of replacement sampling at one full equivalent draw sees about 63% unique
    ↳ points in expectation; subsequent epochs compound coverage.
81 - When surface and volume need different numbers of views, the case is repeated enough times to cover the larger count. The
    ↳ smaller modality returns empty tensors after its own views are covered instead of silently repeating full data.
82 - Validation/test with `--eval-surface-points N --eval-volume-points M`: each case is split into deterministic strided chunks
    ↳ with `torch.arange(view_index, total_points, view_count)`. This is full-fidelity evaluation: every loaded surface point
    ↳ and every loaded volume point is evaluated exactly once, then reaggregated by case.
83 - Set a point limit to `0` only when you intentionally want full-case loading in one batch and have checked memory.
84
    
```

```

85 When launched with `torchrun`, DDP is automatic from `WORLD_SIZE`. Training uses a distributed sampler. Checkpoint-selection
    ↳ validation is exact and distributed across ranks with no padded duplicate eval views, then final `full_val/*` and
    ↳ `test_primary/*` are rerun on rank 0 from the saved checkpoint to match single-model production evaluation semantics.
86
87 ## Model Contract in the Reference train.py
88
89 The reference model interface is:
90
91 ```python
92 out = model(
93     surface_x=batch.surface_x,
94     surface_mask=batch.surface_mask,
95     volume_x=batch.volume_x,
96     volume_mask=batch.volume_mask,
97 )
98 surface_pred_norm = out["surface_preds"] # [B, N_surface, 4]
99 volume_pred_norm = out["volume_preds"] # [B, N_volume, 1]
100 ```
101
102 `batch.surface_mask` and `batch.volume_mask` are required because cases are padded independently. **Do not compute loss or
    ↳ metrics on padding tokens.**
103
104 The reference Transolver backbone must also keep padding masked internally. Slice-attention pooling, residual blocks, final
    ↳ normalization, and output heads should never let padded zero rows contribute to slice tokens or produce nonzero hidden
    ↳ states.
105
106 ## Gradient, Weight, And Slope Telemetry
107
108 Adjust the gradient/weight telemetry defaults to the run length: short debug runs can log more often, while long runs should
    ↳ log less often.
109
110 The W&B stream includes:
111
112 - `train/grad/*` - aggregate gradient health: global norm, RMS, mean absolute gradient, max absolute gradient, zero fraction,
    ↳ non-finite count, parameter norm, and grad-to-parameter norm ratio.
113 - `train/grad_type/<LayerType>/*` - the same statistics grouped by module class, for example `LinearProjection`,
    ↳ `TransolverAttention`, `LayerNorm`, or `TransformerBlock`.
114 - `train/grad_module/<LayerType>/<module_path>/*` - layer-by-layer statistics for every named module with trainable parameters.
115 - `train/grad_param/<LayerType>/<parameter_path>/*` - per-parameter statistics for every trainable tensor.
116 - `train/grad_hist/all` and `train/grad_hist_param/<LayerType>/<parameter_path>` - gradient histograms for distribution drift,
    ↳ spikes, saturation, dead layers, and collapse detection.
117 - `train/weight/*` - aggregate parameter health after each optimizer update: norm, mean, mean absolute value, RMS, standard
    ↳ deviation, min/max, max absolute value, zero fraction, non-finite count, and trainable/frozen tensor counts.
118 - `train/weight_type/<LayerType>/*`, `train/weight_module/<LayerType>/<module_path>/*`, and
    ↳ `train/weight_param/<LayerType>/<parameter_path>/*` - the same parameter statistics grouped by module class, layer path,
    ↳ and parameter tensor.
119 - `train/weight_hist/all` and `train/weight_hist_param/<LayerType>/<parameter_path>` - optional parameter histograms
    ↳ controlled by `--log-weight-histograms`.
120 - `train/grad/global_norm_pre_clip` and `train/grad/clipped` - pre-clip gradient norm and whether the default
    ↳ `--grad-clip-norm 1.0` threshold engaged.
121 - `train/step_skipped`, `train/nonfinite_loss`, and `train/nonfinite_grad` - finite guard diagnostics for poisoned batches.
122 - `train/lr`, `train/base_mse_loss`, `train/surface_loss_weighted`, and `train/volume_loss_weighted` - standard optimization
    ↳ and loss-component diagnostics.
123
124 Keep gradient logging close to `loss.backward()` and before `optimizer.step()` so it represents the update that is about to be
    ↳ applied. Keep weight logging after `optimizer.step()` so it represents the model state after the update. When adding new
    ↳ model blocks, make sure their parameters remain visible through `named_modules()` / `named_parameters()` so the layer-type
    ↳ and layer-path logging stays useful.
125
126 The final metric contract is mandatory. End-of-run `full_val_primary/*` and `test_primary/*` keys must all be present and
    ↳ finite. A run with missing, NaN, or infinite final primary metrics is invalid and should fail loudly rather than being
    ↳ treated as a usable result.
127
128 Slope telemetry is also part of the contract. Every `--slope-log-fraction 0.05` of the estimated optimizer-step budget,
    ↳ `train.py` logs slopes for key curves under `train/slope/*`: losses, grad norms, grad-to-param ratio, RMS gradients, max
    ↳ gradients, MAE curves, and relative-L2 curves. Validation slopes are logged under `val/slope/*` at validation events;
    ↳ final validation and test use `full_val/slope/*` and `test/slope/*` when enough history exists. If you rename metric keys,
    ↳ preserve or document equivalent slope curves.
129
130 Optional early-stop kill thresholds are available through `--kill-thresholds`. The format is a comma- or semicolon-separated
    ↳ list of `STEP:metric<value>`, `STEP:metric<=value>`, `STEP:metric>value`, or `STEP:metric>=value` checks, for example:
131
132 ```
133 --kill-thresholds "500:train/loss<5,2000:val_primary/abupt_axis_mean_rel_l2_pct<25"
134 ```
135
136 Each check is evaluated only when that metric is logged at or after the requested global optimizer step. If the condition
    ↳ fails, the run logs `early_stop/*`, finishes W&B, and skips expensive final validation/test. Use this to discard clearly
    ↳ unstable or hopeless short runs; do not use it to hide final metrics for serious confirmation runs.
137
    
```

```

138 The parser intentionally validates the full format before training starts. Accepted operators are `<`, `<=`, `>`, and `>=`;
139 ↪ steps must be positive integers; values must be finite numbers; metric keys must be spelled exactly as logged.
140
141 ## Metrics
142
143 Checkpoint selection uses an internal scalar mean over the AB-UPT-aligned scalar relative-L2 columns:
144
145 ...
146 abupt_axis_mean_rel_l2_pct = mean(
147     surface_pressure_rel_l2_pct,
148     wall_shear_x_rel_l2_pct,
149     wall_shear_y_rel_l2_pct,
150     wall_shear_z_rel_l2_pct,
151     volume_pressure_rel_l2_pct,
152 )
153 ...
154
155 Primary logged metrics:
156
157 - `val_primary/abupt_axis_mean_rel_l2_pct` - checkpoint selection metric
158 - `full_val_primary/abupt_axis_mean_rel_l2_pct` - best-checkpoint validation metric after training
159 - `test_primary/abupt_axis_mean_rel_l2_pct` - final held-out scalar used for quick run triage
160 - `*_primary/surface_pressure_mae`, `*_primary/wall_shear_mae`, `*_primary/volume_pressure_mae` - target-family MAE diagnostics
161 ↪ target-family relative-L2 diagnostics
162 - `*_primary/surface_pressure_rel_l2_pct`, `*_primary/wall_shear_rel_l2_pct`, `*_primary/volume_pressure_rel_l2_pct` -
163 ↪ target-family relative-L2 diagnostics
164 - `*_primary/wall_shear_x_rel_l2_pct`, `*_primary/wall_shear_y_rel_l2_pct`, `*_primary/wall_shear_z_rel_l2_pct` - per-axis
165 ↪ wall-shear relative-L2 diagnostics
166
167 Lower is better. For paper-facing reporting, use the final `test_primary/*` metrics and include the target-family MAEs plus
168 ↪ `surface_pressure`, vector `wall_shear`, per-axis wall-shear, and `volume_pressure` relative-L2 percentages.
169
170 ### TARGET METRIC ALIGNMENT WITH AB-UPT
171
172 AB-UPT reports relative L2 error in percent, averaged per sample across the split. The paper defines the relative L2 over all
173 ↪ points and output dimensions for a target vector, then averages those per-sample values across test cases.
174
175 For DrivAerML paper-facing reporting, quote the final `test_primary/*` metrics for surface pressure `p_s`, vector wall shear
176 ↪ `tau`, volume pressure `p_v`, and the per-axis wall-shear components when diagnosing residual gaps. This repo logs the
177 ↪ matching columns it can compute:
178
179 - `surface_pressure_rel_l2_pct` maps to AB-UPT `p_s`.
180 - `wall_shear_rel_l2_pct` maps to vector `tau`.
181 - `wall_shear_x_rel_l2_pct`, `wall_shear_y_rel_l2_pct`, `wall_shear_z_rel_l2_pct` map to the per-axis wall-shear columns.
182 - `volume_pressure_rel_l2_pct` maps to AB-UPT `p_v`.
183
184 `abupt_axis_mean_rel_l2_pct` is a convenience aggregate for checkpointing and triage. It is not a standalone AB-UPT benchmark
185 ↪ column, so paper-facing comparisons should quote the individual test fields above.
186
187 ### State-Of-The-Art Targets
188
189 AB-UPT provides aligned public DrivAerML reference targets, including per-axis wall shear. Transolver-3 and Transolver++
190 ↪ provide additional aggregate references for the paper-facing `p_s`, `tau`, and `p_v` comparison. Use the strongest
191 ↪ relevant public target for the reported column, while still reporting this repo's exact held-out `test_primary/*` metrics.
192
193 | Target | This repo metric | AB-UPT reference |
194 |-----|-----|-----|
195 | Surface pressure `p_s` | `test_primary/surface_pressure_rel_l2_pct` | `3.82` |
196 | Vector wall shear `tau` | `test_primary/wall_shear_rel_l2_pct` | `7.29` |
197 | Volume pressure `p_v` | `test_primary/volume_pressure_rel_l2_pct` | `6.08` |
198 | Wall shear `tau_x` | `test_primary/wall_shear_x_rel_l2_pct` | `5.35` |
199 | Wall shear `tau_y` | `test_primary/wall_shear_y_rel_l2_pct` | `3.65` |
200 | Wall shear `tau_z` | `test_primary/wall_shear_z_rel_l2_pct` | `3.63` |
201
202 The per-axis wall-shear table reports `p_s = 3.76` and `p_v = 6.29` in the same AB-UPT/DoMINO comparison setting. Treat the
203 ↪ stricter of the relevant published values as the target band when deciding whether a result is exciting. The camera-ready
204 ↪ SENPAI single-checkpoint result is H147 / PR `#1344` / W&B `k6q4c3on`: `p_s = 3.5634`, `p_v = 3.4014`, vector `tau =
205 ↪ 6.5409`, and AB-UPT-axis mean `5.6648`.
206
207 ### Target Scaling
208
209 The baseline normalizes each target channel with train-split mean/std only. `surface_cp` is already a pressure coefficient,
210 ↪ and the packaged loader does not expose verified per-case freestream or Reynolds metadata for volume pressure or wall
211 ↪ shear. Do not add guessed per-case `p / Re^2` or Cp-style rescaling unless that metadata is explicitly plumbed through and
212 ↪ final validation/test metrics are still computed on the original target units for AB-UPT alignment.
213
214 ## Experiment Length
215
216 Use a mix of shorter and longer experiments:
217
218 - Short runs are for viability: does the idea compile, stay stable, produce sane gradients, and improve early validation
219 ↪ trends?
    
```

```

202 - Longer runs are for confirmation: does a stable idea keep improving once the model has enough update budget to learn surface
    ↪ and volume structure?
203
204 Do not run only short experiments; they can discard ideas before the model has started learning. Do not run only long
    ↪ experiments; they burn the time budget before enough hypotheses are screened. Use the metrics' slope, gradient, and weight
    ↪ logs etc to decide which short runs deserve a longer confirmation run.
205
206 The launch wall-clock limit is a hard cutoff, not a target duration; choose `--epochs` and any step limits according to the
    ↪ evidence needed for the experiment.
207
208 ## Research Workflow
209
210 For substantial architecture or training-strategy changes, preserve main context by using research subagents before
    ↪ implementation. The advisor should deliberately run two complementary streams instead of only local hill-climbing.
211
212 ### Stream 1: Exploit Existing Evidence
213
214 Mine the existing DrivAerML history of experiments before assigning a wave of experiments:
215
216 - Search this target repo's prior branches and PRs, including `main`, `codex/optimized-lineage`, and any previous DrivAerML
    ↪ experiment branches.
217 - Search the PRs from the `wandb/senpai` GitHub repo's `drivaerml-long-20260504` branch which also contains a wealth of
    ↪ previous DrivAerML experiments and analysis.
218 - Assign a subset of students to build on the strongest past ideas rather than rediscovering them: useful preprocessing,
    ↪ target transforms, batching choices, normalization, loss weighting, architecture deltas, and optimizer schedules.
219 - When reusing a past idea, state the source PR/run/branch in the assignment and explain what is being preserved versus
    ↪ changed.
220
221 ### Stream 2: Generate New High-Variance Ideas
222
223 Reserve another subset of students for fresh, aggressive ideas that may not be present in the history. Just focussing on past
    ↪ results above might result in getting stuck in a local optimum. Search broadly across CFD surrogate literature and
    ↪ adjacent ML:
224
225 - DrivAerML and AI-for-CFD work: AB-UPT, UPT, Transolver, Transolver-3, GeoTransolver, DoMINO, GINO/FNO-style operators,
    ↪ graph/mesh transformers, point-cloud networks, geometric encodings, and multi-resolution tokenization.
226 - AI-for-science ideas from physics, chemistry, and biology: equivariant or invariant representations, denoising/pretraining,
    ↪ multiscale message passing, latent neural operators, simulator residual modelling, and uncertainty-aware losses --> there
    ↪ is a lot of interesting work in the broader AI-for-science literature that can be applied to DrivAerML.
227 - Modern transformer and LLM architecture ideas available at run time: Arcee, OLMo, DeepSeek v4, Kimi 2.5 / 2.6, GLM 5.1, and
    ↪ other recent papers for MoE routing, latent/linear attention, token compression, normalization, optimizer, curriculum, and
    ↪ distillation tricks that can plausibly transfer to point/field regression.
228
229 Wild ideas are welcome, but they must still honor the data split, full-fidelity test evaluation, and telemetry contract. A
    ↪ clever idea that cannot be measured on `test_primary/*` is not useful for this sprint.
230
231 ## Roles
232
233 Research is coordinated through GitHub PRs with an advisor/student model. GitHub Issues are used for communication with the
    ↪ human researcher team. See `instructions/prompt-advisor.md` and `instructions/prompt-student.md`.
234
235 ## To Sum up: do everything you possibly can to crush the current reference metrics
236
237 Here are the public reference metrics and the current camera-ready SENPAI single-checkpoint result:
238
239 | Target | This repo metric | H147 single checkpoint | AB-UPT reference |
240 |-----|-----|-----|-----|
241 | Surface pressure `p_s` | `test_primary/surface_pressure_rel_l2_pct` | `3.5634` | `3.82` |
242 | Vector wall shear `tau` | `test_primary/wall_shear_rel_l2_pct` | `6.5409` | `7.29` |
243 | Volume pressure `p_v` | `test_primary/volume_pressure_rel_l2_pct` | `3.4014` | `6.08` |
244 | Wall shear `tau_x` | `test_primary/wall_shear_x_rel_l2_pct` | `5.8155` | `5.35` |
245 | Wall shear `tau_y` | `test_primary/wall_shear_y_rel_l2_pct` | `7.0556` | `3.65` |
246 | Wall shear `tau_z` | `test_primary/wall_shear_z_rel_l2_pct` | `8.4882` | `3.63` |
247
248 Future test metrics should beat H147 while preserving the pressure gains and closing the remaining per-axis wall-shear gap. We
    ↪ have put a massive amount of effort into this research programme to get this far, now is the time for you to step up and
    ↪ deliver the results we need.

```

K.5. Local TandemFoil2 Program Contract

```

1 # TandemFoil2 - program context for senpai agents
2
3 This is the active SENPAI problem package: a clean Transolver-based implementation seeded from the `tcapelle/kagent`
    ↪ TandemFoil competition. See `KAGENT_SOURCE.md` for provenance.
4
5 ## Task
6

```

```

7 Predict the 3D airflow velocity field around F1 front wings (tandem-foil geometries) given mesh-node features. The
  ↳ TandemFoilSet has seven pickle files spanning `raceCar` (land-inverted) and `cruise` (aerial-freestream) domains with
  ↳ structured holdouts on front-foil camber and Reynolds number (see `organizer/SPLITS.md`).
8
9 ## Data contract
10
11 - Input `x`: `(N, 24)` float32 per mesh node - geometric + flow features (see `data.py`). **Do not edit `data.py`.**
12 - Target `y`: `(N, 3)` float32 - velocity components `(u, v, w)`.
13 - Mask `is_surface`: `(N,)` bool - mesh nodes on the airfoil surface. Velocity is zero there (no-slip).
14 - `N` varies per sample, up to ~300K points. Each sample is ~100K 3D points on average.
15 - Splits materialised under `/mnt/<pvc-mount>/datasets/tandemfoil/splits_v2/{train, val_single_in_dist, val_geom_camber_rc,
  ↳ val_geom_camber_cruise, val_re_rand}/` as individual `.pt` files. Runner: `organizer/prepare_splits.py` (one-off, reads
  ↳ pickles + manifest, writes .pt shards).
16
17 ## Primary metric
18
19 - **`val/l2_error`** - mean L2 velocity error across all four val splits, equal-weighted. Lower is better.
20 - Tracked per split: `val/<split_name>/l2_error`.
21 - Ranking metric for the leaderboard.
22
23 ## Constraints
24
25 - **96GB VRAM ceiling** per student pod. 100K-point samples are large - subsample, chunk, or use memory-efficient attention.
26 - **Timeout** is wall-clock, controlled by `SENPAI_TIMEOUT_MINUTES` (default 30 min for smoke runs). Honor it - training
  ↳ should checkpoint and exit cleanly.
27 - **Max epochs** via `SENPAI_MAX_EPOCHS` (default 50; current run uses 999). Early-stop on validation plateau.
28
29 ## Files you may edit
30
31 - `train.py` - the model, optimizer, training loop, logging. Primary target of experimentation.
32 - `predict.py` - inference / prediction assembly. Edit when you change model I/O shape.
33 - `EXPERIMENT_JOURNAL.md` - append one entry per attempted experiment, success or failure. Failed experiments are the most
  ↳ valuable.
34
35 ## Files that are read-only
36
37 - `data.py` - data loader contract. Changing this breaks comparability across students.
38 - `organizer/*` - split manifest, scoring, and materialisation are the ground-truth contract.
39 - `KAGENT_SOURCE.md` - attribution record.
40 - `README.md`, `program.md`, `KAGGLER_AGENT.md` - description and rules.
41
42 ## Git workflow (critical)
43
44 This repo is pulled into the SENPAI runner as a **git submodule** under `target/tandemfoil2/`. All your commits, branches, and
  ↳ PRs live in THIS repo (`morganmcg1/tandemfoil2`) on branch `kagent_royal_rumble` - **not in `wandb/senpai`**. Your
  ↳ entrypoint script sets this up automatically; when in doubt verify `git remote -v` points at `morganmcg1/tandemfoil2.git`.
45
46 - Base branch: `kagent_royal_rumble`.
47 - Your working branch: `${RESEARCH_TAG}/student-${STUDENT_NAME}` (set by the entrypoint).
48 - PR target: `gh pr create --repo morganmcg1/tandemfoil2 --base kagent_royal_rumble`.
49 - Never `git commit` in the parent senpai repo - that repo is read-only from the agent's perspective.
    
```

K.6. kagent Comparator Prompt

```

1 You are an autonomous kaggler in a live competition against other coding agents. Your goal: **predict 3D airflow velocity
  ↳ fields around F1 front wings better than everyone else.**
2
3 **BEFORE WRITING ANY CODE: read `README.md` and `EXPERIMENT_JOURNAL.md` completely.** It describes the data format, model
  ↳ contract, metrics, and memory constraints.
4
5 # Current Leaderboard
6
7 !`cat /mnt/new-pvc/predictions/${RESEARCH_TAG}/leaderboard.md`
8
9 ## Key files
10
11 - `README.md` - competition description, data format, metrics, rules. **Read cover to cover before starting.**
12 - `data.py` - data loader. **Read-only.**
13 - `train.py` - training template. Fill in your model where it says `NotImplementedError`.
14 - `predict.py` - prediction template. Same: fill in your model loading code.
15
16 ## The experiment loop
17
18 You work on branch `${RESEARCH_TAG}/kaggler/<your-name>`. It's already checked out.
19
20 LOOP FOREVER:
21
    
```

```

22 1. **Check the competition.** Read the leaderboard: `cat /mnt/new-pvc/predictions/$RESEARCH_TAG/leaderboard.md`. Query W&B for
    ↳ the best runs. Know where you stand.
23 2. **Formulate a hypothesis.** What will you try next?
24 3. **Modify `train.py`** (and `predict.py` if needed). Do **not** edit the journal yet - the entry lands after you know the
    ↳ outcome.
25 4. **Commit the code change only.** `git add train.py predict.py && git commit -m "<what you're trying>"`. Keep the journal
    ↳ out of this commit so a later reset doesn't erase it.
26 5. **Run training**: `python train.py --agent <your-name> --wandb_name "<your-name>/<description>" > run.log 2>&1`
27     - Read results: `grep "Best:" run.log` and `tail -5 run.log`
28     - If error: `tail -50 run.log` for the traceback.
29 6. **Run predictions** (if training succeeded): `python predict.py --checkpoint <path> --agent <your-name> > pred.log 2>&1`
30 7. **Keep or discard the code change:**
31     - If improved → stage the best checkpoint alongside it:
32       `git add checkpoints/best.pt && git commit -m "ckpt: val/l2=<score>"`.
33     - If worse or crashed → reset the code commit: `git reset --hard HEAD-1`.
34
35     The best checkpoint is always mirrored to `checkpoints/best.pt` (local git path) and to
    ↳ `/mnt/new-pvc/kagent/$RESEARCH_TAG/$KAGGLER_NAME/checkpoints/model-<run_id>/checkpoint.pt` (PVC, durable).
36 8. **Always update the journal - even for failures.** After step 7 completes (kept **or** discarded), append an entry to
    ↳ `EXPERIMENT_JOURNAL.md` covering hypothesis, change, result, verdict, notes. Then commit and push the journal on its own
    ↳ so the record of what you tried survives regardless of whether the code landed:
37
38     ```
39     git add EXPERIMENT_JOURNAL.md && git commit -m "journal: <short summary>" && git push
40     ```
41
42     Failed experiments are the most valuable entries to keep - they prevent you (and future iterations) from repeating the same
    ↳ dead end. Never skip this step.
43 ## Key challenges
44
45     - **Memory**: 100k 3D points per sample. You MUST address this - subsample points, use efficient architectures, or both.
46     - **Turbulence**: The hard part is high-frequency turbulent components. The laminar flow is easy (just copy the input).
47     - **No-slip BC**: Velocity is zero on the airfoil surface (`idcs_airfoil`). Enforce this as a constraint.
48
49 ## Metrics
50
51 **Primary**: `val/l2_error` - mean L2 velocity error. Lower is better. This is what the leaderboard ranks by.
52
53 ## Constraints
54
55     - `data.py` is read-only
56     - Training timeout: controlled by `MAX_TIMEOUT_MIN` env var
57     - VRAM: 96GB. Don't OOM - this dataset is large.
58
59 ## NEVER STOP
60
61 Once the loop begins, do NOT pause to ask the human. You are autonomous. If you run out of ideas, think harder - check what
    ↳ the leaders are doing, search the web for new approaches. The loop runs until the human interrupts you.
62
63 ## WIN
64
65 Your objective is to top the leaderboard. If you are stuck with a low scoring solution, don't be afraid to try radical
    ↳ changes, marginal improvements on your low scoring solution are not going to cut it!

```